

Operating System (OS)

CS232

Persistence: I/O Devices

Dr. Muhammad Mobeen Movania

Mr. Abid Butt

Outlines

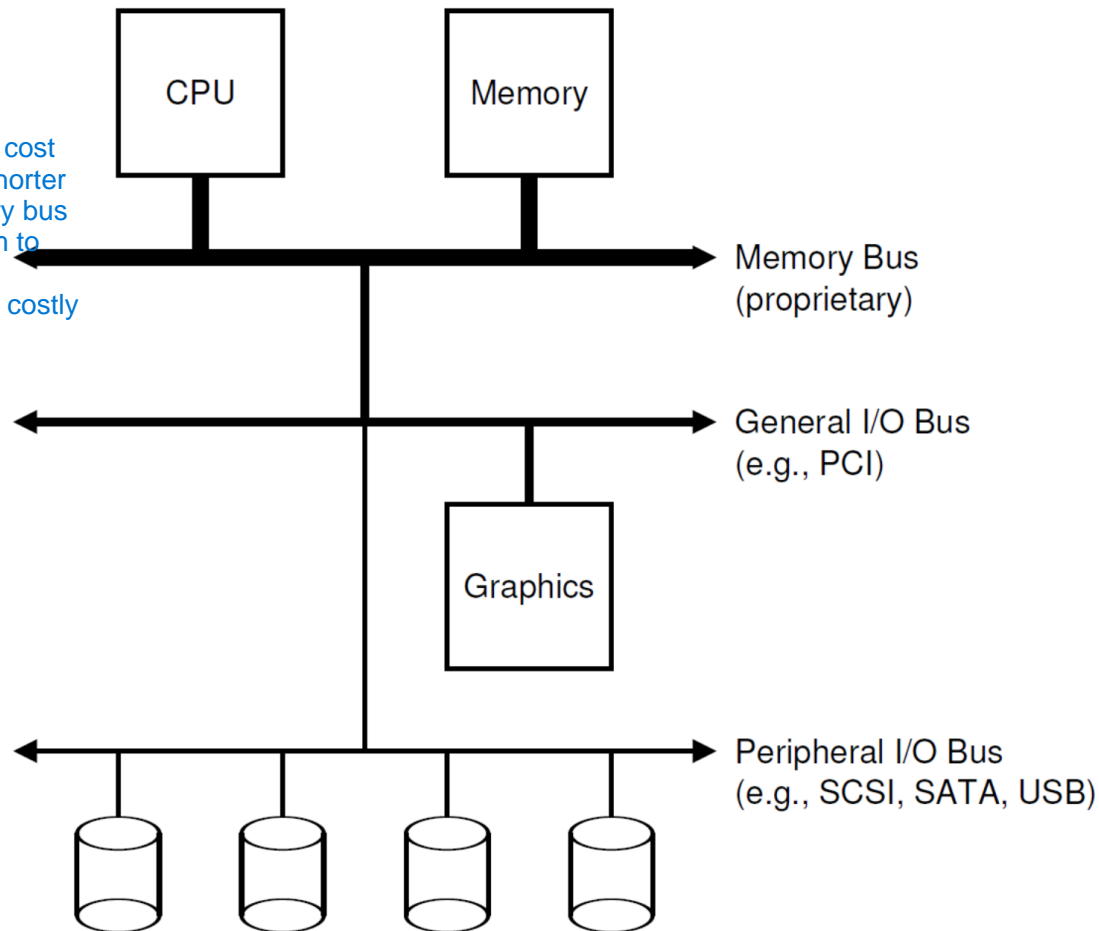
- Significance of I/O
- I/O interfaces in a prototypical and modern system
- Typical I/O device
- Typical protocol
- Communicating with devices
- Summary

Significance of I/O

- No serious program can be considered complete without I/O
- Its important to vary processing of data based on inputs and outputs
- Challenge is to support a wide range of I/O products in as general an abstraction as possible

Simple Prototypical System

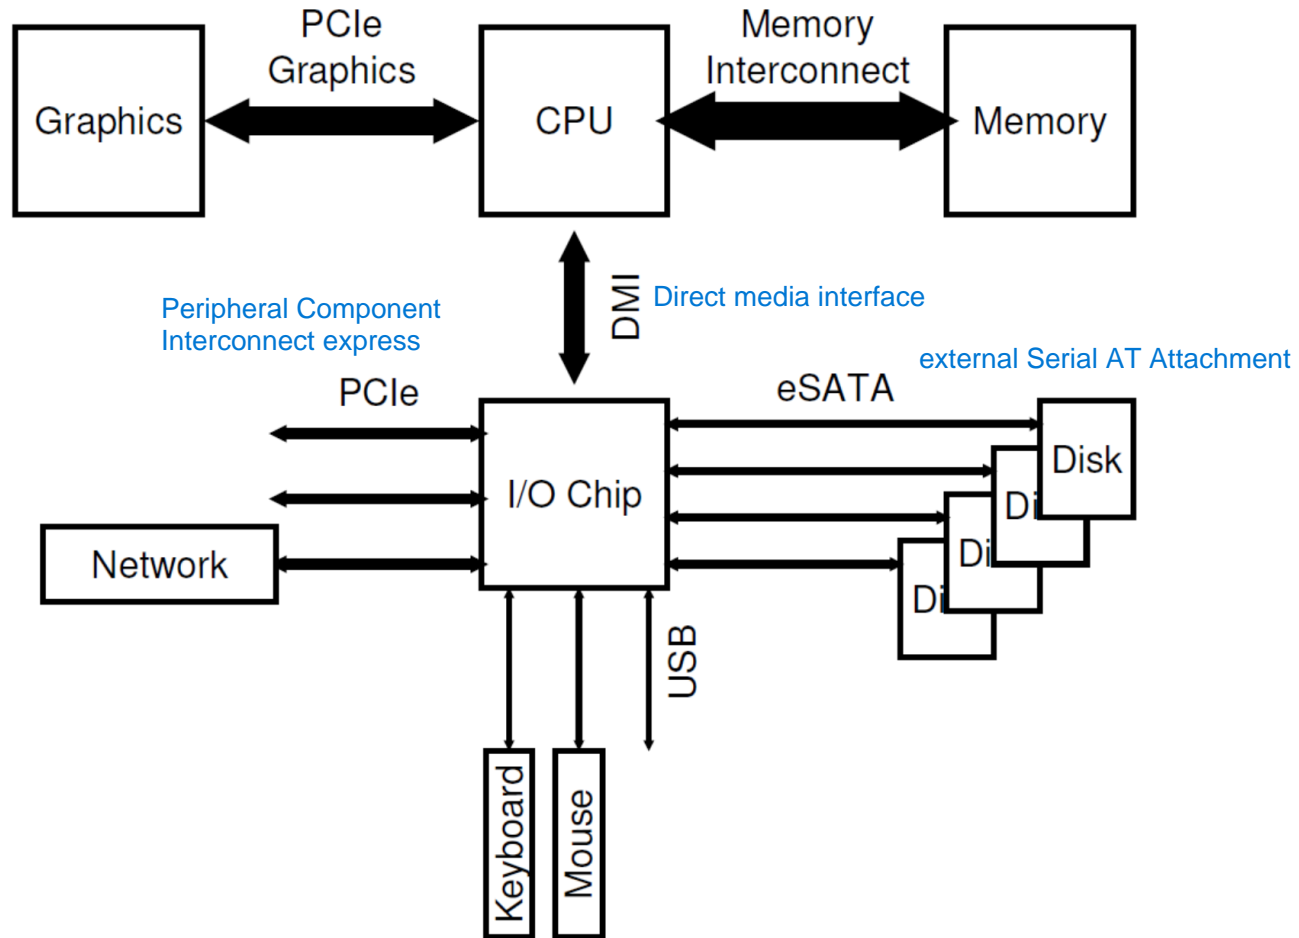
- why hierarchical: physics and cost
- physics: faster bus must be shorter
 - high-performance memory bus does not have much room to plug devices
- cost: high-performance bus is costly



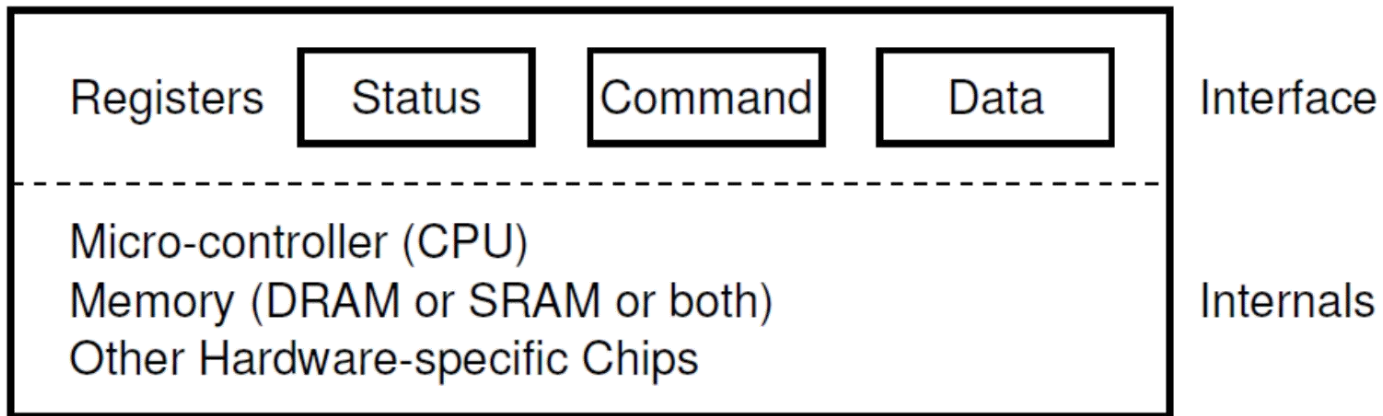
disks, mice, keyboards (slower devices)

Typical I/O Interfaces in a modern system

- specialized chipsets
- faster point-to-point interconnects



Typical I/O Device



Three types of I/O

When main CPU is involved with data movement

- Programmed I/O (aka Polling)
- Interrupt driven I/O
- DMA

Typical Protocol

- Uses polling

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (Doing so starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

Positive aspect:

- Simple and working

Negative:


- inefficiencies (waiting wastes CPU time, could switch to something else)
- inconveniences

Interrupts

- Let computation and I/O overlap
 - OS writes the data and command
 - Puts the calling process to sleep and shifts to another task
 - After finishing the task, the device raises an interrupt
 - The ISR is called
 - The OS reads the result of the I/O, wakes the sleeping process

Not always the best solution:


- Imagine a device that performs its tasks very quickly
 - If fast, best to poll
 - If slow, use interrupts
 - If unknown speed: use hybrid (poll for a while, then use interrupt)
- Networks:
 - huge stream of incoming packets
 - each generating interrupt
 - can cause "livelock"
- Coalescing:
 - device to wait before issuing interrupt
 - can bunch together other requests that got done



Direct Memory Access (DMA)

Data movement via ...


- Programmed I/O is also very expensive!



- DMA engine can transfer data b/w RAM and device w/o CPU intervention

OS would program the DMA engine by telling it:

- where the data lives in memory
- how much data to copy
- where to send it to



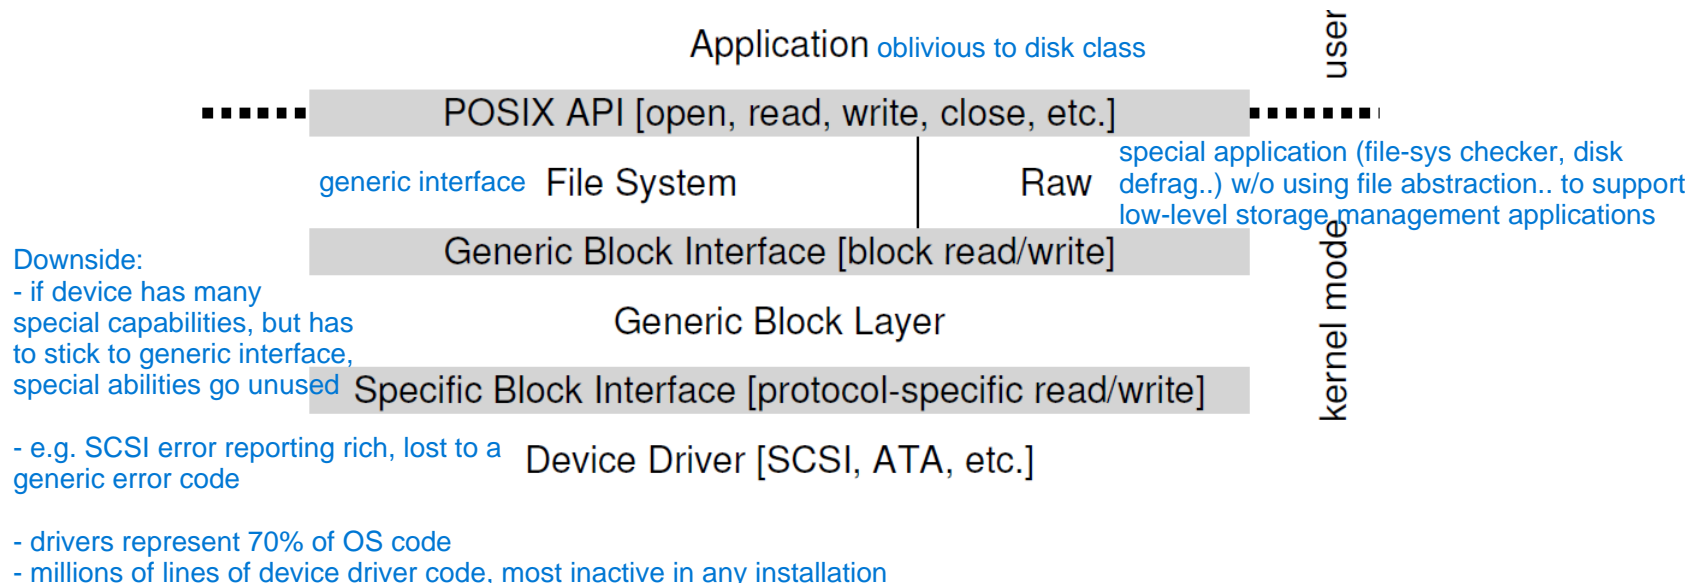
Communicating with devices

- Two main techniques:
 - Explicit I/O instructions
 - Intel has in and out instructions
 - Memory mapped I/O
- usage:
out reg, port
- privileged: only OS has access
- hardware makes device regs available as if they were memory locations
- OS issues a load or store for reading/writing the address
- hardware routes load/store to device instead of memory

OS is oblivious to device details

- Each device has its own characteristics
- How do we do it?
 - Abstraction!
- Device driver is a layer that hides the details

e.g. file system to work on top of SCSI, IDE, USD, and so forth
file system should be oblivious to all the details



Summary

- We have a basic understanding of how OS manages I/O devices
- Three types of I/O
 - Programmed I/O
 - interrupt driven I/O and
 - DMA
- Two approaches to accessing device registers
 - explicit I/O instructions and
 - memory-mapped I/O,
- Device driver has been presented, showing how the OS itself encapsulates low-level details and makes it easier to build the OS in a device-neutral fashion