

Operating System (OS)

CS232

Persistence: Hard disk drives

Dr. Muhammad Mobeen Movania

Outlines


- What is a HDD
- Physical geometry of HDD
- HDD working
- I/O time and disk scheduling
- Disk head positioning algorithms
- Disk scheduling issues
- Summary

Hard Disk Drive (HDD)

- An I/O device
- Used for persistent storage
 - main form persistent storage for decades
 - file system technology predicated on their behavior
- Consists of sectors (512 byte each) numbered 0 to $n-1$
- Writes to a sector are atomic
 - many file systems will read or write 4KB at a time
 - but only write to a sector is guaranteed to be atomic
- Bigger read/writes are possible but not atomic
- Accessing blocks near to each other is faster than accessing blocks far away from each other

HDD – physical geometry


- Platter
 - Made of metal (aluminum) coated with magnetic material on both sides
 - Data is written by inducing magnetic changes
 - A disk may have one or more platters
- Surface: a side of a platter
- Spindle: binds together multiple platters
- Speed: Revolutions per minute (RPM) - 7.2-15k
- Tracks: concentric circle made up of sectors
 - 1000s on a single surface
 - tightly packed (100s fit into width of a hair)
- A Disk Head: is used to read/write data
 - one head per surface
- A Disk Arm: is used to move the head, position it over the right track



HDD – working

- building a model one track at a time
- assume a simple disk with a single track

- Rotational Delay
 - To read track 0?
 - Max rotational delay?




- if R is delay of full rotation
- $R/2$ to wait for 0
- worst case is sector 5

HDD – working: seek operation

- Seek phases = acceleration + coasting + deceleration + settling

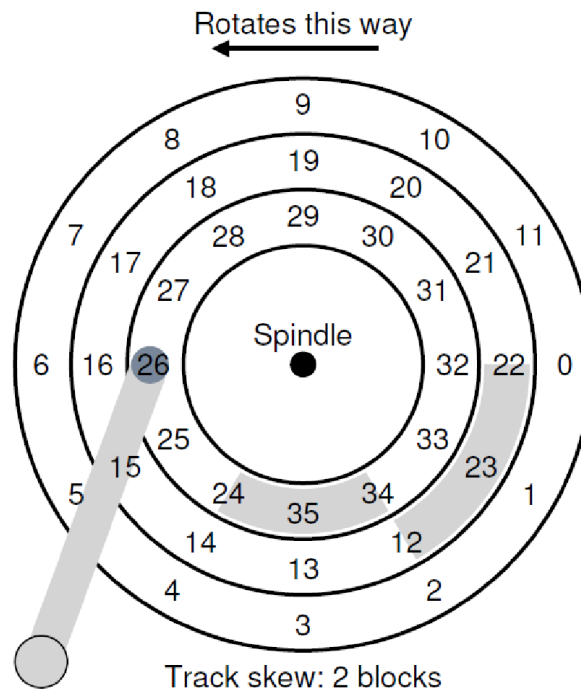
-significant (0.5-2ms) as the drive must be certain to find the right track



- I/O time = seek time + rotate time + transfer

Track skew

- To facilitate transfers crossing track boundaries



- to prevent one full rotational delay when head moves to neighboring track

Track buffer

- Works as a Cache
 - Reads and stores multiple consecutive sectors for future use
- Writing can have two strategies:
 - Write back - acknowledging write completed when written in cache
 - Write through -acknowledging write completed when actually written to disk

HDD specs

	high performance SCSI drive Cheetah 15K.5	built for capacity Barracuda
Capacity	300 GB	slow but pack as many bytes as possible 1 TB
RPM	engineered to spin as fast as possible 15,000	7,200
Average Seek	low seek time 4 ms	9 ms
Max Transfer	transfer data quickly 125 MB/s	105 MB/s
Platters	4	4
Cache	16 MB	16/32 MB
Connects via	SCSI	SATA

I/O Time

for a "random" load of 4 KB

- small reads to random locations on
- the disk common in database mgmt

$$T_{I/O} = T_{seek} + T_{rotation} + T_{transfer}$$

$$R_{I/O} = \frac{Size_{Transfer}}{T_{I/O}}$$

On the Cheetah:

this is avg seek time.

full seek will be ~2-3 times longer

$$T_{seek} = 4 \text{ ms}, T_{rotation} = \frac{0.5 * (60 * 1000 / 1500)}{2} \text{ ms} = 2 \text{ ms}, T_{transfer} = \frac{4 \text{ KB} / 125 \text{ MB/s}}{1000} = 30 \text{ microsecs}$$

vanishingly small

$$T_{I/O} = 6 \text{ ms}, R_{I/O} = \frac{0.004 \text{ MB} / 6e-3 \text{ s}}{1} = 0.66 \text{ MB/sec}$$

$$\text{For barracuda: } T_{I/O} = 13.2 \text{ ms}, R_{I/O} = 0.31 \text{ MB/sec}$$

$$9 + 0.5 * (60 * 1000 / 7200) + (4000 / 105000000) * 1000$$

$$0.004 / 13.2e-3$$

Sequential Transfer (100MB)

reads large #sectors consecutively from the disk w/o jumping around

- Cheeta T I/O = 800 ms
100MB / 125MB/s
- Barracuda T I/O = 950 ms
100MB / 105MB/s

	Cheetah	Barracuda
$R_{I/O}$ Random	0.66 MB/s	0.31 MB/s
$R_{I/O}$ Sequential	125 MB/s	105 MB/s

- Conclusion: Drive performance varies heavily whether we are doing random small transfers or big sequential transfers !!

Disk scheduling

Module of OS that decides which I/O request to the disk should be scheduled next


- Multiple I/O requests of varying characteristics
- How can we perform them efficiently?
- Disk scheduler tries to schedule them to reduce average delay.
- Length of job is known ??
- SJF == SSTF == NBF
- Starvation?

shortest job first

shortest seek time first

nearest block first

- when drive geometry is not available to OS
- OS instead sees an array of blocks



SSTF can lead to starvation

- a steady stream of requests to the inner track
- requests to any other tracks would be ignored

- SCAN algorithm remedies this
- It moves the head in a sweep (from inner to outer tracks and vice versa)
- The next request served will be the closest one but not on the same track!

- freezes the queue to be serviced when it is doing a sweep


- circular scan. sweep in one direction only, then reset

- Variations: FSCAN, CSCAN
- Also called the Elevator algorithm.
 - imagine when going floor 10 -> 1, someone got in at 3 and pressed 4 (closer than 1)

SPTF (shortest positioning time first)

- Which sector to read next: 16 or 8?
- SSTF fine if seek-time >>> rotation time
- SSTF and SCAN do not take into account the time to rotate.
- Problem: often OS does not know the internal details of the disk!

- so it's usually performed inside a drive



Disk Scheduling – issues

- Where is the scheduling performed?
 - OS - happened in older systems
 - HDD - modern systems can implement SPTF accurately
- I/O merging
 - Imagine 3 requests to read from sectors 33, 8, 34 respectively
 - scheduler should merge 33 and 34 into a single request
 - further reordering performed on the merged request
 - important to do at OS level, reduces number of requests sent to the disk
- How long the system should wait before issuing an I/O request to disk?
 - Immediately (work conserving), or
 - After some time (non work conserving)
 - by waiting a new and "better" request may arrive at the disk, and thus overall efficiency is increased

Summary

- We saw how disk drives work
- We saw different strategies that are used to track the head positioning on tracks and sectors
- OS and disk drive manufacturers try to design a tradeoff for an affordable and performant persistent disk storage