

# Operating System (OS)

## CS232

Persistence: Files and Directories

Dr. Muhammad Mobeen Movania

# Outlines

- Files and directories
- File system interface
- Creating, accessing and deleting files
- Sequential and random I/O
- Shared file table entries
- Hard and symbolic links
- Permission bits and access control list (ACL)
- Summary

# Files and Directories

- We saw two virtualizations earlier
  - Process (virtualization of CPU)
  - Address space (virtualization of memory)
- We will now look at virtualization of storage
- Two main abstractions
  - File
  - Directory

# Files and Directories (2)

- File
  - A linear array of bytes
  - Has a low-level name (which is often a number) called the **inode number**
  - Often files have an extension but its not mandatory to have contents as per extension
- Directory
  - Similar to a file, has a low-level name (inode number)
  - Contains a list of tuples containing (user-readable name, low-level name)
  - Each entry is either a file or another directory
  - Directory within directory allow to create a **directory tree (or directory hierarchy)**




Figure 39.1: An Example Directory Tree

# File System Interface

- Set of API for creating, accessing and deleting files
- Underneath they may call one or more system calls
- All internal implementation details are hidden from the user

# Creating Files

- Uses open system call with O\_CREATE flag

```
int fd = open("foo",  
              O_CREAT|O_WRONLY|O_TRUNC,  
              S_IRUSR|S_IWUSR) ;
```

- The second parameter is a set of flags
  - (O\_CREAT) creates the file if it does not exist,
  - (O\_WRONLY) ensures that the file can only be written to, and,
  - (O\_TRUNC) truncates file (if the file already exists) to a size of zero bytes thus removing any existing content.
- The third parameter specifies permissions,
  - (S\_IRUSR|S\_IWUSR) file is readable and writable by the owner
- The return value is a file descriptor that is used to access file in later operations
- File descriptors are stored per-process in **proc** struct

# Reading and Writing Files

- We can use **strace** tool to extract system calls made by programs
- Use **fsync (fd)** call to force write flushing

```
prompt> strace cat foo
...
open("foo", O_RDONLY|O_LARGEFILE)      = 3
read(3, "hello\n", 4096)                = 6
write(1, "hello\n", 6)                  = 6
hello
read(3, "", 4096)                       = 0
close(3)                                = 0
...
prompt>
```

# Non-sequential read/write

- Use the lseek system call to reposition the read/write pointers in file

```
off_t lseek(int file_des, //file descriptor
            off_t offset, //offset in bytes
            int whence);  //from start (SEEK_SET)
                        //from current (SEEK_CUR)
                        //from end (SEEK_END)
```

System Calls	Return Code	Current Offset
fd = open("file", O_RDONLY);	3	0
read(fd, buffer, 100);	100	100
read(fd, buffer, 100);	100	200
read(fd, buffer, 100);	100	300
read(fd, buffer, 100);	0	300
close(fd);	0	–



# Reading from a single file using two descriptors

System Calls	Return Code	OFT[10] Current Offset	OFT[11] Current Offset
<code>fd1 = open("file", O_RDONLY);</code>	3	0	–
<code>fd2 = open("file", O_RDONLY);</code>	4	0	0
<code>read(fd1, buffer1, 100);</code>	100	100	0
<code>read(fd2, buffer2, 100);</code>	100	100	100
<code>close(fd1);</code>	0	–	100
<code>close(fd2);</code>	0	–	–

# Reading after lseek

System Calls	Return Code	Current Offset
<code>fd = open("file", O_RDONLY);</code>	3	0
<code>lseek(fd, 200, SEEK_SET);</code>	200	200
<code>read(fd, buffer, 50);</code>	50	250
<code>close(fd);</code>	0	–

# Shared File Table Entries

```
int main(int argc, char *argv[]) {
    int fd = open("file.txt", O_RDONLY);
    assert(fd >= 0);
    int rc = fork();
    if (rc == 0) {
        rc = lseek(fd, 10, SEEK_SET);
        printf("child: offset %d\n", rc);
    } else if (rc > 0) {
        (void) wait(NULL);
        printf("parent: offset %d\n",
               (int) lseek(fd, 0, SEEK_CUR));
    }
    return 0;
}
```

Figure 39.2: Shared Parent/Child File Table Entries (**fork-seek.c**)

- Output

```
prompt> ./fork-seek
child: offset 10
parent: offset 10
prompt>
```

# Shared File Table Entries




Figure 39.3: Processes Sharing An Open File Table Entry

# Renaming a file

- fsync call flushes writes to ensure persistence
- Rename is achieved through
  - mv foo bar
- strace call reveals

```
int fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC,  
              S_IRUSR|S_IWUSR);  
write(fd, buffer, size); // write out new version of file  
fsync(fd);  
close(fd);  
rename("foo.txt.tmp", "foo.txt");
```

# Deleting a file and directory

- Delete file
  - Use rm command
  - Using strace, we see that rm calls unlink

```
prompt> strace rm foo
...
unlink("foo")                                = 0
...
```

- Delete directory
  - Use rmdir command
  - Directory must be empty or rmdir fails silently

# Reading a Directory

- We may simulate the ls command

```
int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");
    assert(dp != NULL);
    struct dirent *d;
    while ((d = readdir(dp)) != NULL) {
        printf("%lu %s\n", (unsigned long) d->d_ino,
              d->d_name);
    }
    closedir(dp);
    return 0;
}
```

```
struct dirent {
    char          d_name[256]; // filename
    ino_t         d_ino;       // inode number
    off_t         d_off;       // offset to the next dirent
    unsigned short d_reclen;    // length of this record
    unsigned char  d_type;      // type of file
};
```

# Hard links

- `link()` function
  - makes an entry in the file system tree
  - takes two arguments, an old pathname and a new one;
- When you “link” a new file name to an old one, you essentially create another way to refer to the same file.
- Use the `ln` command as follows

```
prompt> echo hello > file
prompt> cat file
hello
prompt> ln file file2
prompt> cat file2
hello
```



# rm command and Unlink

- After creating a hard link to a file, there is no difference between the original file name (file) and the newly created file name (file2)
- They are both just links to the underlying metadata about the file, which is found in inode number 67158084.

```
prompt> ls -i file file2
67158084 file
67158084 file2
prompt>
```

- To remove a file from the file system, we call unlink().

```
prompt> rm file
removed 'file'
prompt> cat file2
hello
```

# rm command and Unlink (2)

```
prompt> echo hello > file
prompt> stat file
... Inode: 67158084      Links: 1 ...
prompt> ln file file2
prompt> stat file
... Inode: 67158084      Links: 2 ...
prompt> stat file2
... Inode: 67158084      Links: 2 ...
prompt> ln file2 file3
prompt> stat file
... Inode: 67158084      Links: 3 ...
prompt> rm file
prompt> stat file2
... Inode: 67158084      Links: 2 ...
prompt> rm file2
prompt> stat file3
... Inode: 67158084      Links: 1 ...
prompt> rm file3
```

# Symbolic links

- Hard links have a few restrictions
  - you can't hard link to a directory (for fear it will create a cycle in the directory tree)
  - you can't hard link to files in other disk partitions (because inode numbers are only unique within a particular file system)
- Soft link are created with the -s flag to ln

```
prompt> echo hello > file
prompt> ln -s file file2
prompt> cat file2
hello
```

```
prompt> stat file
... regular file ...
prompt> stat file2
... symbolic link ...
```

```
prompt> ls -al
drwxr-x---  2 remzi remzi   29 May  3 19:10 ./
drwxr-x--- 27 remzi remzi 4096 May  3 15:14 ../
-rw-r----- 1 remzi remzi    6 May  3 19:10 file
lrwxrwxrwx  1 remzi remzi    4 May  3 19:10 file2 -> file
```

# Symbolic links

- Soft links might generate dangling pointers

```
prompt> echo hello > file
prompt> ln -s file file2
prompt> cat file2
hello
prompt> rm file
prompt> cat file2
cat: file2: No such file or directory
```

# Permission bits and access control list (ACL)

- On unix based systems, we can use `ls -l` command to see the file permission bits

```
prompt> ls -l foo.txt  
-rw-r--r--  1 remzi wheel  0 Aug 24 16:29 foo.txt
```

- The first character (-) here just shows the type of the file (- for a regular file), d for a directory, l for a symbolic link)
- Next nine bits are for permission
  - The permissions consist of three groupings: what the **owner** of the file can do to it, what someone in a **group** can do to the file, and finally, what anyone (sometimes referred to as **other**) can do

# Summary

- We saw how file and directories are implemented
- We saw how the OS system calls abstract the low level details during file I/O
- The file system interface in UNIX systems (and indeed, in any system) is seemingly quite rudimentary, but there is a lot to understand