

Operating System (OS)

CS232

Process

Dr. Muhammad Mobeen Movania

Outlines

- What is a process?
- Multitasking vs Multithreading
- Process abstraction
- Process API
- Process creation
- Process states and state transitions
- Process data structure
- Summary

What is a process?

- A running instance of a program
- Process abstraction
 - allows operating system to run many programs simultaneously like a browser and word process running at the same time
 - Makes system easy to use
- How does OS give illusion of many CPUs to allow so many programs to run at the same time?
 - OS virtualizes the CPU

Illusion of an infinite supply of CPU

- At any instant, there can be 100 or 1000 processes running at the same time on an OS
- OS creates this illusion by virtualizing the CPU
- Time Sharing of CPU
 - OS runs one process, then stops it and then runs another process (context switching) and so on

Multitasking vs Multithreading

- Multitasking is also call time slicing or time sharing of CPU
 - Idea is to rapidly switch processes to give an illusion of many processes running at the same time
 - At any time, only one process is running on the CPU
- Multithreading
 - Idea is to run several concurrent processes at the same time
 - CPU must have more than one core to truly appreciate the benefits
 - At any time, more than one processes are running on the CPU

Single Threaded Multitasking (Single Core CPU)

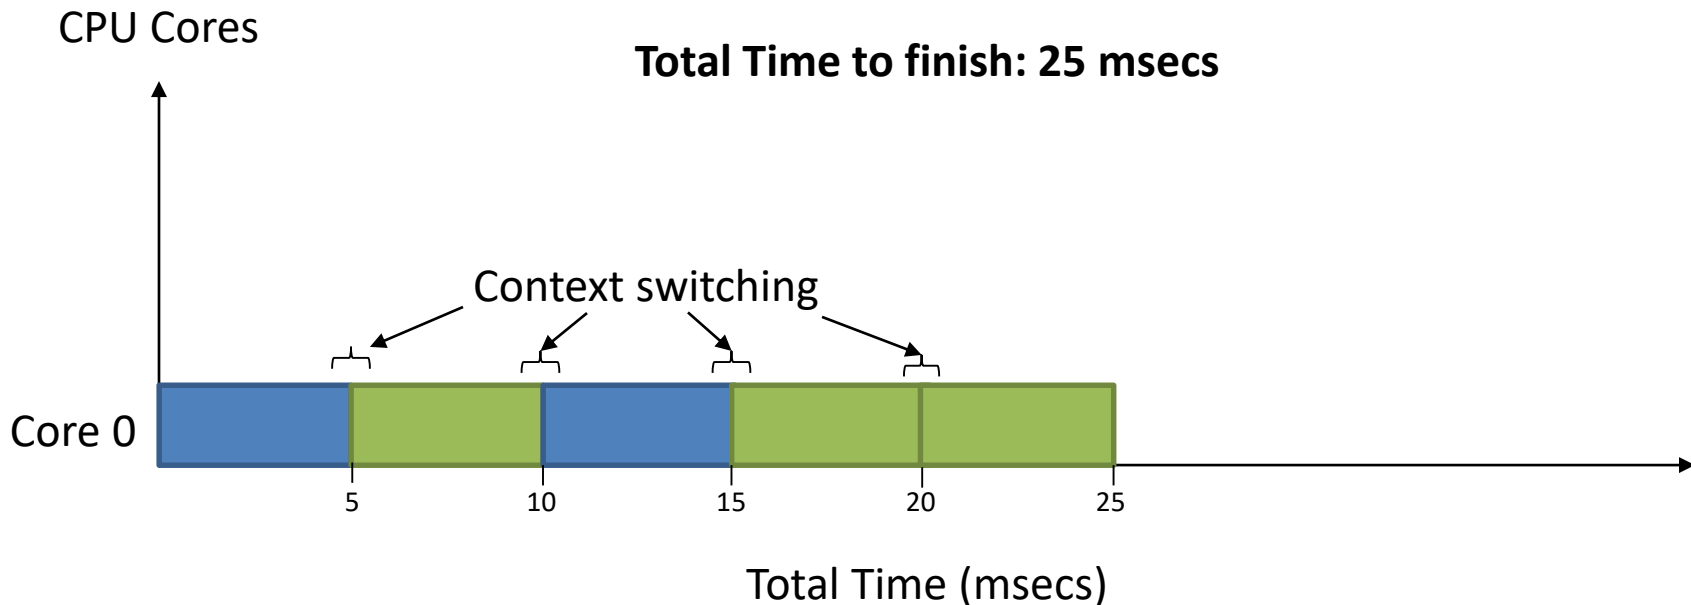
Num of CPU Cores : 1

Time slice: 5 msecs

P1 (10 msecs)



P2 (15 msecs)



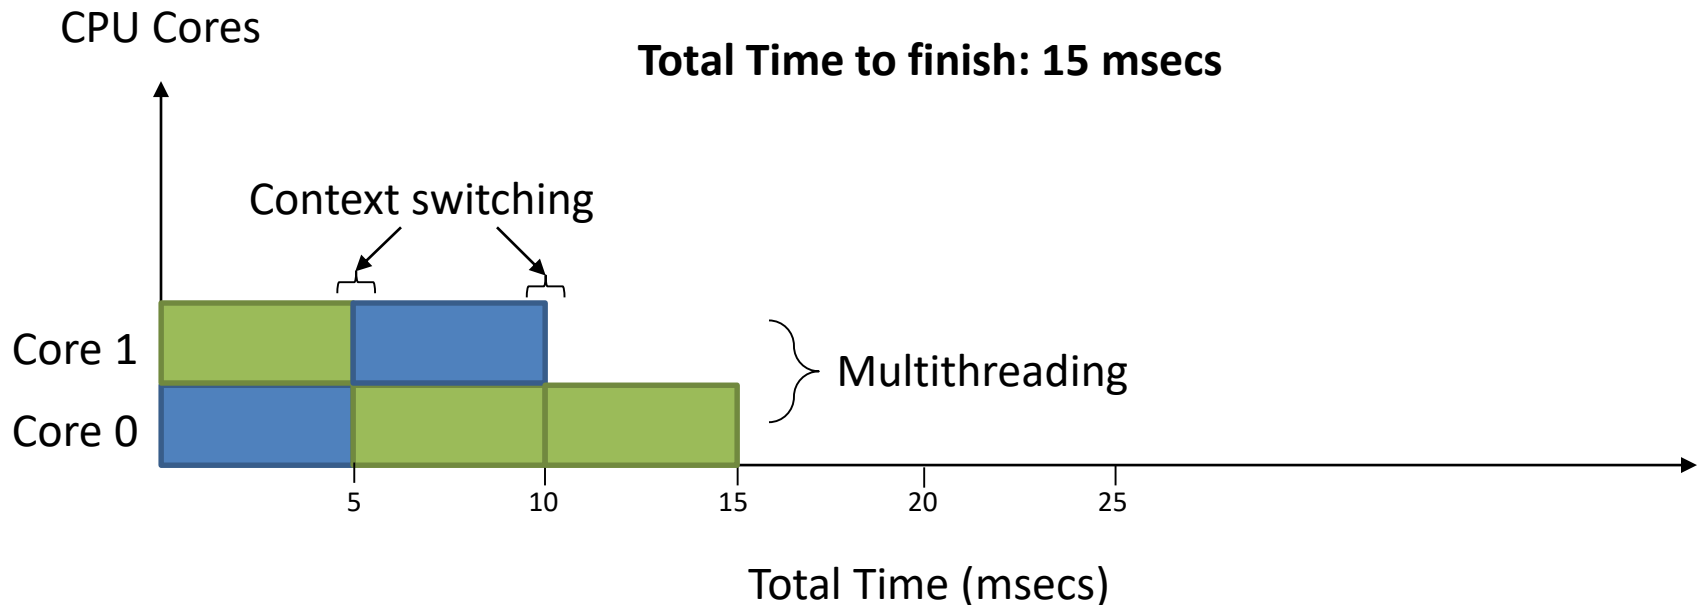
Multithreaded Multitasking (Dual Core CPU)

Num of CPU Cores : 2
Time slice: 5 msecs

P1 (10 msecs)



P2 (15 msecs)



Process Abstraction

- Has to store
 - process or machine state (values of CPU registers, PC/IP, stack pointer (SP), frame pointer(FP) etc.)
 - memory the process can address (address space)
 - I/O information (list of open files)

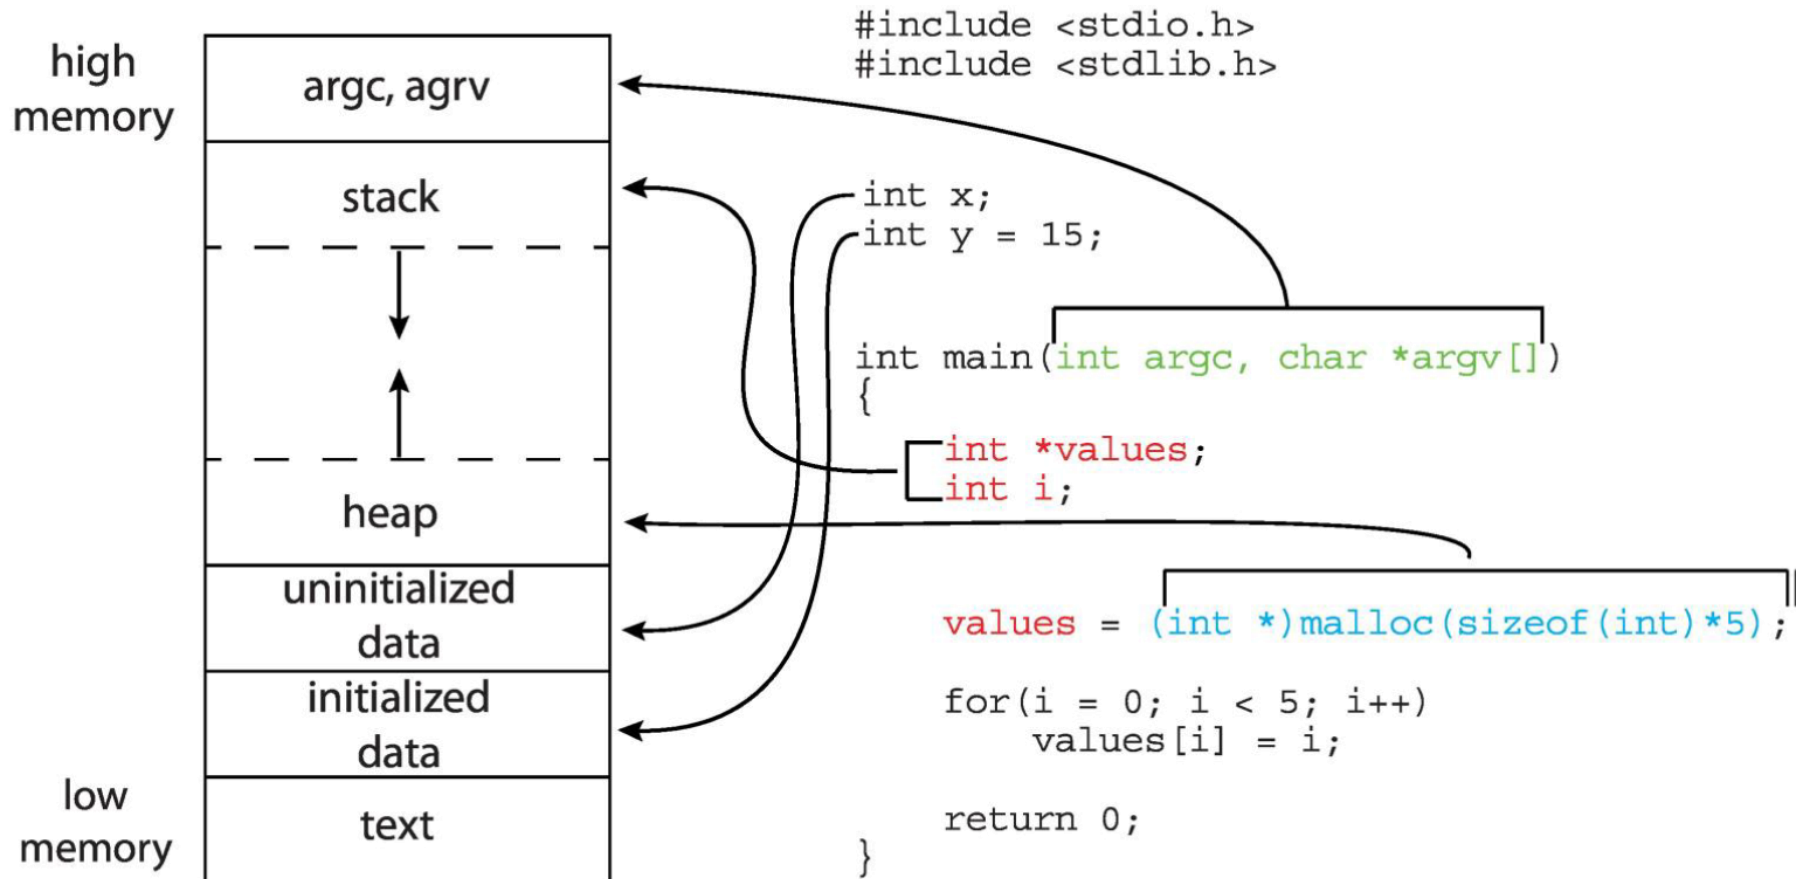
Process API

- OS must provide API to
 - Create
 - Destroy
 - Wait
 - Misc. Control (suspension of a running process)
 - Status (time and state of process)

Process Creation

- What happens when you double click an exe?
 - OS loads code and static data into the address space of process in RAM
 - Loading can be eager (all at once) or lazy (on demand)
 - OS allocates some memory for process's run-time stack (for storing function local variables and return addresses, command line args) and some for heap (dynamic memory allocations)
 - OS also initializes I/O related interfaces
 - OS finally transfers control to the process main function

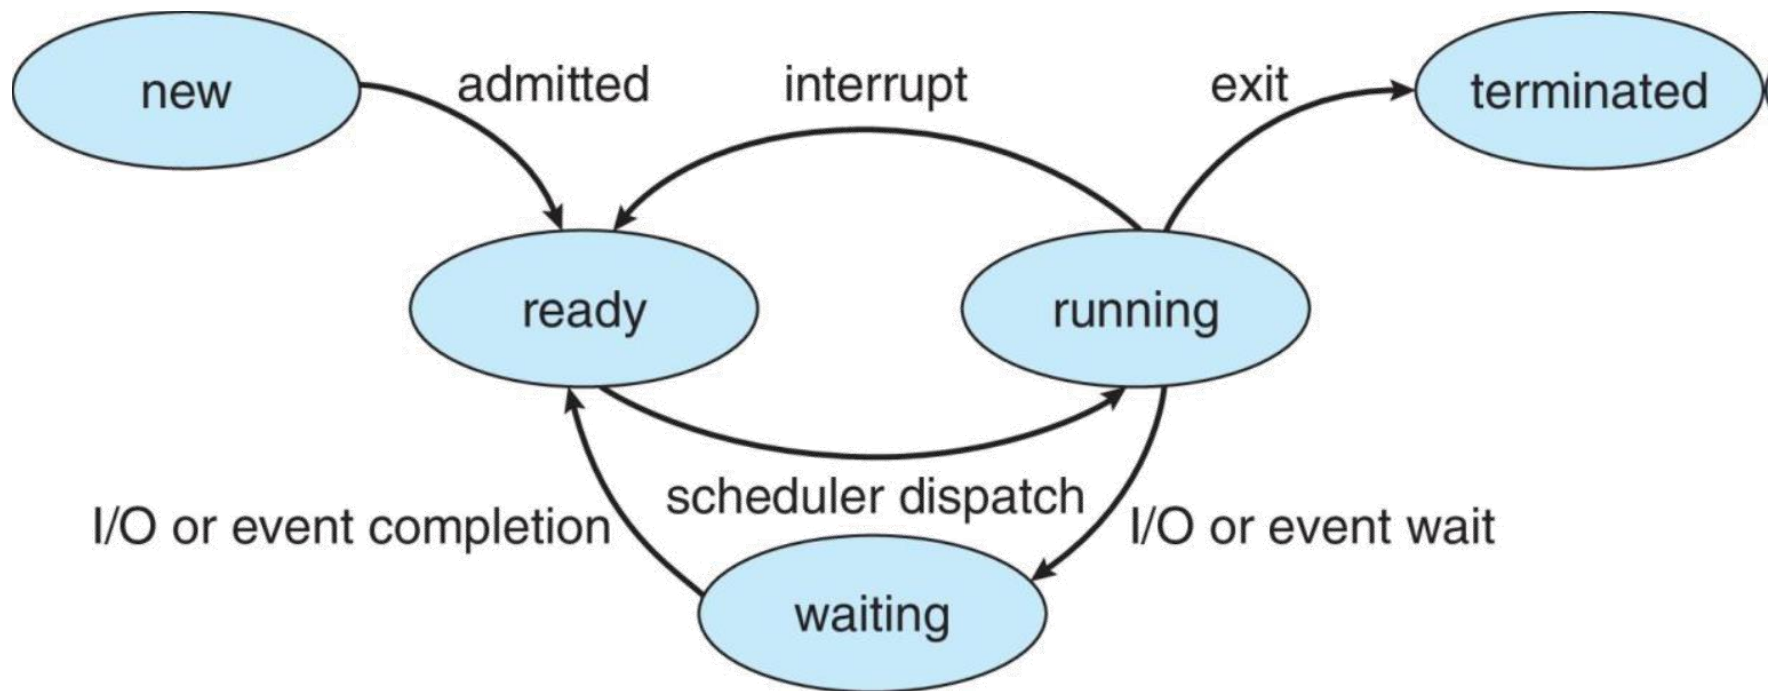
Process Creation



Process States

- A process can be in one of the five states
 - New (Process is created)
 - Running (Process is executing instructions)
 - Ready (Process is ready but OS has not chosen it)
 - Waiting/Blocked (Process has performed some operation like I/O request so it cannot continue until that operation finishes)
 - Terminated (Process has finished execution)

Process State Transition



Examples of Process Transitions without I/O

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	—	Running	
6	—	Running	
7	—	Running	
8	—	Running	Process ₁ now done

Figure 4.3: Tracing Process State: CPU Only

Examples of Process Transitions with I/O

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked,
5	Blocked	Running	so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

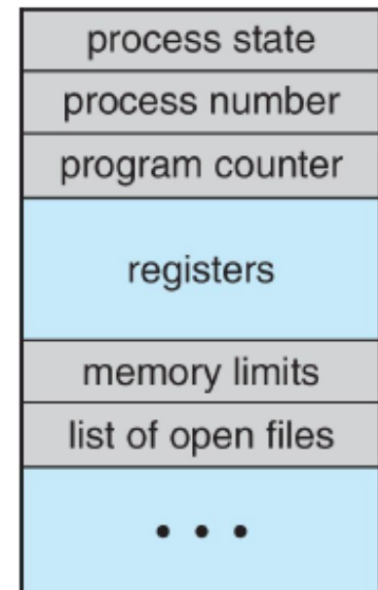
Figure 4.4: Tracing Process State: CPU and I/O

Process Data Structure

- Process list contains information of all processes in system
- Each entry in the process list is a PCB

Information associated with each process(also called **task control block**)

- Process state – running, waiting, etc.
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files



PCB - Implementation

- Register context holds contents of registers for the stopped process so that execution can continue when the process is scheduled again

```
struct context {  
    int eip;    //insn ptr  
    int esp;    //stack ptr  
    int eax;    //gp, originally accumulator  
    int ebx;    //gp  
    int ecx;    //gp, originally counter  
    int edx;    //gp  
    int esi;    //gp, for string ops is source index  
    int edi;    //gp, for string ops is dest. index  
    int ebp;    // base ptr  
}; //segment registers: CS, DS, SS, ES
```

PCB – Implementation (2)

```
// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                              // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If !zero, sleeping on chan
    int killed;               // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;        // Current directory
    struct context context;   // Switch here to run process
    struct trapframe *tf;    // Trap frame for the
                              // current interrupt
};
```

Summary

- Process is a major **OS abstraction** of a **running program**
- A **process list** contains information about all processes in the system. Each entry is found in what is sometimes called a **process control block (PCB)**
- Process state is stored in **PCB** which includes
 - Contents of memory in process address space
 - Values of CPU registers (PC/IP,SP,FP)
 - I/O information (open files list)
- Process goes through **5 states** during its lifetime and transitions among them due to events (scheduled or descheduled, or waiting for I/O)