

# Operating System (OS)

## CS232

Process: API and Implementation

Dr. Muhammad Mobeen Movania

# Outlines

- What is an API?
- Dual mode operation and transitions
- POSIX
- Types of system calls
- Process management system calls
- `fork()`, `wait()`, `exec()` with examples
- Summary

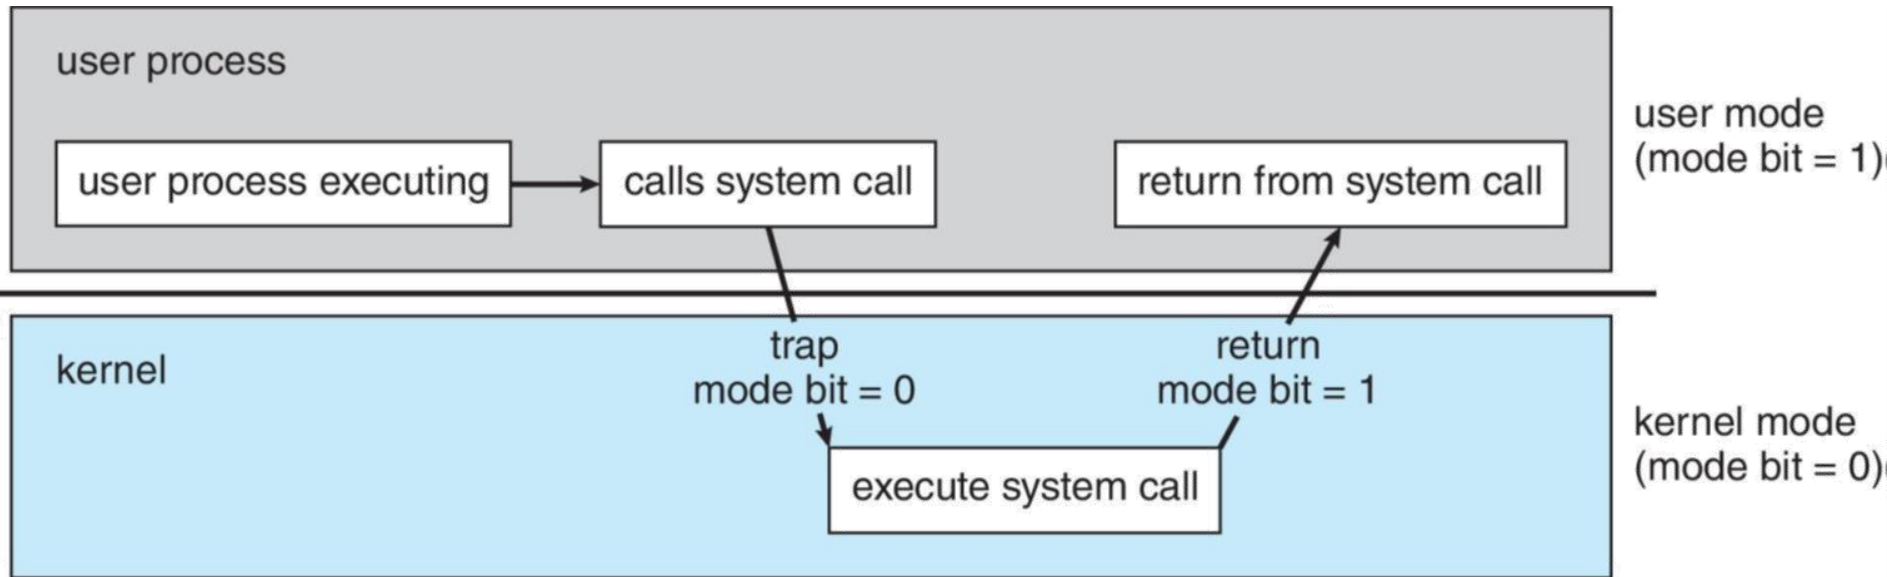
# What is an API?

- API (Application programming interface) is a collection of functions that are provided to users to control any system or program
- All OS provide functions called **system calls**
- System calls
  - Provide access to hardware and other privileged accesses to user processes
  - Are always run in kernel mode (privilege mode)

# Dual mode operation

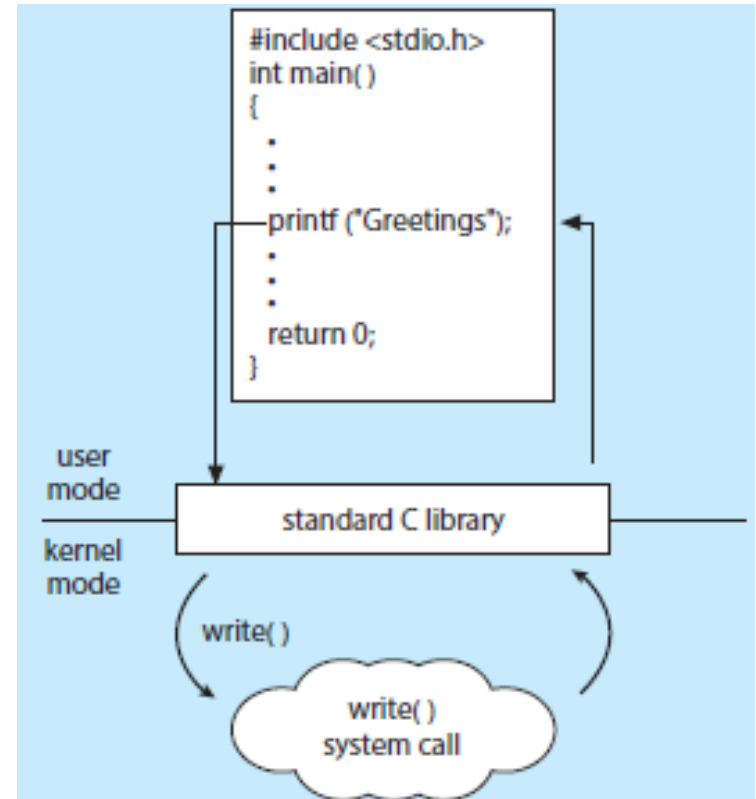
- OS provide two modes of operation
  - User mode (non-privileged)
  - Kernel mode (privileged)
- Why
  - Protection of hardware and other system components
- How
  - A special mode bit is provided in hardware
  - Changed through a system call
- When running a system call, the process must transition from user mode to kernel mode

# Transition from user to kernel mode



# POSIX

- POSIX (Portable OS Interface) Standard
  - API Standard
  - Ensures compatibility across different OS
  - Programs using POSIX API sure to run on POSIX-compliant OS
  - Most OS provide some sort of POSIX compliance
- Libraries provide an easy-to-use interface to make system calls
  - C language has libc library
  - `printf()` calls the `write()` system call



# Types of System Calls

- System calls may be grouped into the following 6 types
  - Process control
  - File management
  - Device management
  - Information maintenance
  - Communications
  - Protection

# Examples of Windows and Unix System Calls

## EXAMPLES OF WINDOWS AND UNIX SYSTEM CALLS

The following illustrates various equivalent system calls for Windows and UNIX operating systems.

|                                | Windows   | Unix                                   |
|--------------------------------|---|--|
| <b>Process control</b>         | CreateProcess()<br>ExitProcess()<br>WaitForSingleObject()                           | fork()<br>exit()<br>wait()             |
| <b>File management</b>         | CreateFile()<br>ReadFile()<br>WriteFile()<br>CloseHandle()                          | open()<br>read()<br>write()<br>close() |
| <b>Device management</b>       | SetConsoleMode()<br>ReadConsole()<br>WriteConsole()                                 | ioctl()<br>read()<br>write()           |
| <b>Information maintenance</b> | GetCurrentProcessID()<br>SetTimer()<br>Sleep()                                      | getpid()<br>alarm()<br>sleep()         |
| <b>Communications</b>          | CreatePipe()<br>CreateFileMapping()<br>MapViewOfFile()                              | pipe()<br>shm_open()<br>mmap()         |
| <b>Protection</b>              | SetFileSecurity()<br>InitializeSecurityDescriptor()<br>SetSecurityDescriptorGroup() | chmod()<br>umask()<br>chown()          |



# Process Management System Calls in Unix based Systems

- The following 4 functions are provided for process management
  - `fork()` : for creation of a new process
  - `exec()` : for creation of a new process
  - `exit()` : for termination of a process
  - `wait()` : to wait for a created process to complete

# fork()

- Creates a new process which is an exact copy of its parent process
- On success, the fork() system call returns twice
  - For the newly created child process, in which case it returns 0
  - For the parent process, in which case it returns the PID of the child process
- Execution of program continues to the statements after fork() with each process having its own address space

# Example of fork()

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int)
                                                       getpid());
    }
    return 0;
}
```

## Output

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```


# What's going on?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0)
    {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    }
    else if (rc == 0)
    {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    }
    else
    {
        // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```

parent process  
(PID: 29146)

child process  
(PID: 29147)



# Issues in the last code?

- Non-deterministic
  - After the fork call, the child process or parent process might run its statements depending on who gets scheduled on the CPU

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

# wait()

- Sometimes, its useful for a parent process to wait for a child process to finish.
- Achieved through wait() or waitpid().
- Parent process calls wait() to delay its execution until the child finishes.
- When child is done, wait() returns to the parent
- Why add wait()
  - Makes output **deterministic** that is you are always sure that the output statements of child process will be printed first before the parent's output statements are printed

# Example of wait()

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (main)
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n", rc, rc_wait, (int)
            getpid());
    }
    return 0;
}
```

## Output

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (rc_wait:29267) (pid:29266)
prompt>
```

# What's going on?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();

    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (main)
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n", rc,
              rc_wait, (int) getpid());
    }
    return 0;
}
```


## Output

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (rc_wait:29267) (pid:29266)
prompt>
```

parent process  
(PID: 29266)

child process  
(PID: 29267)

parent process waits for  
child process to finish





# exec()

- Used when you want to run a program which is different from the calling program
- Linux has six variants of exec()
  - execl, execlp(), execle(), execv(), execvp(), and execvpe()
- The example code on the next slide runs a word counting program (wc) with the source file given as argument
  - wc returns no. of lines, words and bytes in the given file

# Example of exec()

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>
```

```
int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else {
        // parent goes down this path (main)
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n", rc, rc_wait, (int)
            getpid());
    }
    return 0;
}
```

## Output

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
      29      107      1030 p3.c
hello, I am parent of 29384 (rc_wait:29384) (pid:29383)
prompt>
```

# What's going on?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>
```

```
int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else {
        // parent goes down this path (main)
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n", rc,
              rc_wait, (int) getpid());
    }
    return 0;
}
```

parent process  
(PID: 29383)

child process  
(PID: 29384)

parent process waits for child  
process to finish



# More about exec()

- `exec()` does not create a new process, it transform the currently running process (p3) into a different program (wc)
- How?
  - From the given executable name and arguments, it loads code (and static data) from the executable and overwrites current process's code (and static data)
  - Stack, heap and other parts of memory space are reinitialized
- `exec()` if successful never returns
- Combining `fork()` and `exec()` allows creation of Shells, output redirection (>) etc.
- Unix pipes (|) are implemented by combination of `fork()` with `pipe()` system call

# Summary

- We looked at three process creation API, `fork()`, `exec()` and `wait()`
  - `fork()` is used to create a new child process which is exact replica of the parent process
  - `exec()` allows a child to execute an entirely new program
  - `wait()` allows a parent to wait for its child to complete execution
- Unix shell uses `fork()`, `wait()`, and `exec()` to launch user commands
- Separation of `fork()` and `exec()` enables features like input/output redirection, pipes, and other cool features