

# Operating System (OS)

## CS232

Scheduling Algorithm: Multi-Level Feedback Queue (MLFQ)

Dr. Muhammad Mobeen Movania

# Outlines

- Recap
- Scheduling Algorithm: MLFQ
- Basic rules of MLFQ
- How MLFQ changes priority?
- How MLFQ avoids starvation and gaming?
- MLFQ parameter tuning
- Summary

# Recap

- Ideally, we want a scheduler that minimizes turnaround time (like SJF, STCF) and also minimizes response time (like RR)
- The Problem is
  - The scheduler has no prior knowledge about the process like how much time it will need to finish and its general behavior
- One solution is
  - observe the process's behavior over time and adjust scheduler properties dynamically
- This is the core concept on which MLFQ are based

# Scheduling Algorithm: MLFQ

- MLFQ has a number of distinct **queues**, each assigned a different **priority level**.
- At any given time, a job that is ready to run is on a single queue.
- MLFQ uses **priorities** to decide which job should run at a given time: a job with higher priority (i.e., a job on a higher queue) is chosen to run.
- For jobs with same priority, round robin scheduling is used.

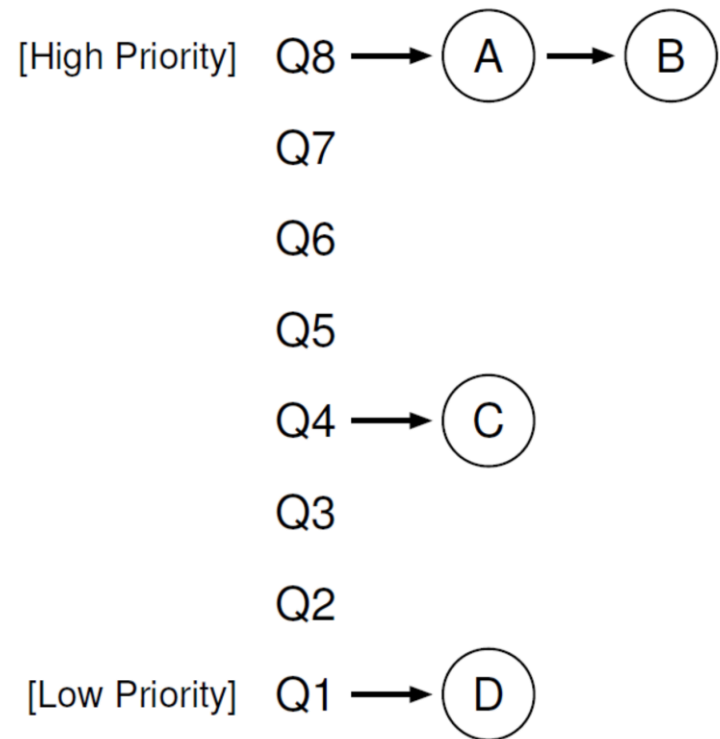


Figure 8.1: MLFQ Example

# MLFQ

- Key Idea:
  - *Vary priority* of a job on its *observed behavior*
- Priority of a process which releases CPU repeatedly (for e.g. due to I/O) is high.
- Priority of a process which uses CPU intensively for long duration is low.
- MLFQ tries to *learn* about processes as they run, and thus use the *history* of the job to predict its *future behavior*

# Basic Rules for MLFQ

- **Rule 1:** If  $\text{Priority}(A) > \text{Priority}(B)$ , A runs (B doesn't).
- **Rule 2:** If  $\text{Priority}(A) = \text{Priority}(B)$ , A & B run in RR.
- **Rule 3:** When a job enters the system, it is placed at the highest priority (the topmost queue).
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue).
- **Rule 5:** After some time period  $S$ , move all the jobs in the system to the topmost queue. (priority boosting)

# How does MLFQ change priority?

- When a job enters the system, it is given highest priority (Rule 3).
- When a job consumes its time share fully its priority is reduced no matter if its consumed fully in one go or in pieces (Rule 4).
- Example
  - Single long running job

# Example 1 – Single Long-Running Job

- Time slice: 10 ms
- Job enter highest priority queue (Q2)
- If time slice expires, the job's priority is reduced and its shifted to the next lower level queue (Q1)
- If time slice expires again, the job's priority is reduced and its shifted to the lowest level queue (Q0)




Figure 8.2: Long-running Job Over Time



# Example 2 – A short job comes in

- Time slice: 10 ms
- Each process moves to next lower priority queue when it finishes its time slice
- Two jobs in system
  - A (CPU bound) (Black)
  - B (Interactive) (Gray)
- Process B arrives at  $T=100$ , runtime=20ms
- Process A moves to  $Q_0$  as its runtime is long whereas B finishes in 20ms so it is moved till  $Q_1$




Figure 8.3: Along Came An Interactive Job

# Key Takeaway

- MLFQ first assumes that the given process might be a short process giving it the highest priority.
  - If it is indeed a short process it finishes execution quickly
  - If it is a long process, it moves down to next lower queue giving it a lower priority
  - This way MLFQ approximates SJF

# Example 3 – CPU+I/O intensive process

- Time slice: 10 ms
- Two jobs in system
  - A (CPU bound) (Black)
  - B (Interactive) (Gray)




Figure 8.4: A Mixed I/O-intensive and CPU-intensive Workload

- Process B takes CPU for 1 ms so it remains on a higher priority queue (Q2)

# Issues with running many interactive jobs

- CPU bound processes are starved as they are never scheduled
- Scheduler can be gamed by a process that issues an I/O request relinquishing the CPU thus keeping the process at the highest priority
- Solution
  - Instead of allowing process's time share to reset upon an I/O request, the total time share taken by a process with and without I/O should be considered (Rule 4)

# Avoiding starvation

- Periodically boost priority of all the jobs
  - Prevents CPU bound processes to be starved as they are guaranteed CPU share after a given time period (S) when they are moved to the topmost queue.




Figure 8.5: Without (Left) and With (Right) Priority Boost

# Avoid gaming due to I/O

- Use better accounting that is keep note of the time share a process has used irrespective of whether it finished its time slice as one long burst or many small bursts.




Figure 8.6: Without (Left) and With (Right) Gaming Tolerance

# MLFQ Parameter Tuning

- Few key parameters
  - How many queues should there be?
  - How big should the time slice be per queue?
  - How often should priority be boosted?
- No easy answer
  - Usually OS's provide config files with default values that OS admins can modify as per need instead of using **Voo-Doo Constants**
  - Solaris MLFQ provides a table that admins can change
    - 60 queues
    - Time slices: 20ms – x00ms
    - Priorities boosted every 1 sec
  - FreeBSD scheduler calculates priority as a function of CPU usage
    - CPU usage is decayed over time resulting in giving less priority

# Summary

- We discussed MLFQ
  - Multiple levels of queues
  - Uses feedback to determine priority depending on the process behavior without prior knowledge of job's running time
  - Uses five rules when assigning priorities to and scheduling processes
  - Provides the best of both worlds by
    - Minimizing turnaround times like (SJF, STCF)
    - Minimizing response time like (RR)