

Operating System (OS)

CS232

Scheduling Algorithm: Proportional Share Scheduling

Dr. Muhammad Mobeen Movania

Outlines

- Why Proportional Share Scheduler
- Types of Proportional Share Schedulers
 - Lottery Scheduling/Fair Share Scheduler
 - Stride Scheduling/Deterministic Fair Share Scheduler
 - Completely Fair Scheduler (CFS)
- Weighing Parameter
- Implementation Details
- Summary

Why Proportional Share Schedulers?

- Key Idea
 - Instead of optimizing for turnaround or response time, a scheduler might instead try to guarantee that each job obtain a certain percentage of CPU time
- Key Questions
 - How to share the CPU proportionally?
 - What are the key mechanisms?
 - How effective are they?

Types of Proportional Share Schedulers

- Three types on the mechanism through which the proportional share of CPU is given
 - Lottery Scheduling (Hold a lottery randomly to see which process runs next)
 - Stride Scheduling (Deterministic lottery scheduling based on a stride value which is inversely proportional to the no. of tickets)
 - CFS Scheduler (Fairly divide CPU among all processes through accumulating virtual runtime)

Lottery Scheduling

- Each process is awarded a certain number of ***tickets***.
- Key Idea:
 - The ***percentage of tickets*** a process has represents its ***share of the resource*** in question
- Mechanism:
 - The Lottery scheduler holds a lottery randomly to pick a winning ticket
 - The process to which the winning ticket belongs to runs next

Example - Lottery Scheduling

- Total Tickets: 100 (A 75% share and B 25% share)
 - Process **A** (Tickets: 0-74), Process **B** (Tickets:75-99)
- Example of Lottery scheduler's winning tickets

63 85 70 39 76 17 29 41 36 39 10 99 68 83 63 62 43 0 49

- Resulting schedule

A **A** **A** **A** **A** **A** **A** **A** **A** **A** **A** **A** **A** **A** **A**
 B **B** **B** **B**

- Summary
 - No guarantees due to randomness
 - We get 16/20 (80% A) and 4/20 (20% B)
 - More chances and we are likely to get the desired percentages

Ticket Currency

- Mechanism to manipulate tickets in lottery scheduling
- Ticket Currency
 - allows a user with a set of tickets to allocate tickets among their own jobs in whatever currency they would like
 - the system then automatically converts said currency into the correct global value

Example - Ticket Currency

- Two Users A and B each with 100 tickets
 - User A
 - Runs two jobs A1 and A2 gives each 500 tickets in A's currency (equal to 50 tickets in global currency)
 - User B
 - Runs a single job and gives it 10 tickets (equal to 50 tickets in global currency)
- Lottery is held on global ticket currency to determine which job runs
- Why use currency
 - Allows ticket transfer: a process can hand off its tickets to another process (useful in client/server setting)
 - Allows ticket inflation: a process can temporarily raise or lower the number of tickets it owns

Lottery - Implementation

- Easy to implement as we just need
 - A good random number generator
 - A data structure to keep track of processes
 - Total number of tickets

```
1 // counter: used to track if we've found the winner yet
2 int counter = 0;
3
4 // winner: use some call to a random number generator to
5 //           get a value, between 0 and the total # of tic
6 int winner = getrandom(0, totaltickets);
7
8 // current: use this to walk through the list of jobs
9 node_t *current = head;
10 while (current) {
11     counter = counter + current->tickets;
12     if (counter > winner)
13         break; // found the winner
14     current = current->next;
15 }
16 // 'current' is the winner: schedule it...
```



Issues with Lottery Scheduling

- 1) Unfairness
 - Two processes A and B with the same number of tickets (100), each job should finish at the same time but this is generally not the case
 - Unfairness Metric (U) = $\frac{\text{completion_time}_A}{\text{completion_time}_B}$
 - U approaches 1 only as the jobs run for a significant number of time slices
- 2) How to assign tickets?
 - Open problem
 - We assume the user knows which ticket should be assigned to which job
- 3) Non-deterministic
 - Based on randomness which approximates to correct answer when outcomes are more




Figure 9.2: Lottery Fairness Study

Stride Scheduling – Deterministic Proportional Share Scheduling

- Stride Scheduling
 - $stride \propto \frac{1}{Ticket_share}$
- Each process has a counter (*pass* value)
- At any given time scheduler will:
 - Select the process with minimum pass value
 - Increment its pass value by its stride
 - Schedule it

Example – Stride Scheduler

- Three processes, A(100), B(50), C(250) tickets
- Total tickets = 10000
 - $stride_A = \frac{10000}{100} = 100$,
 - $stride_B = \frac{10000}{50} = 200$,
 - $stride_C = \frac{10000}{250} = 40$
- Each process starts with a pass value of 0

Example – Stride Scheduler

Pass(A) (stride=100)	Pass(B) (stride=200)	Pass(C) (stride=40)	Who Runs?
0	0	0	A
100	0	0	B
100	200	0	C
100	200	40	C
100	200	80	C
100	200	120	A
200	200	120	C
200	200	160	C
200	200	200	...

Figure 9.3: **Stride Scheduling: A Trace**

Issues - Stride Scheduler

- Example Summary
 - C ran five times, A twice, and B just once, exactly in proportion to their ticket values of 250, 100, and 50
- Lottery scheduling achieves the proportions probabilistically over time but stride scheduling gets them exactly right at the end of each scheduling cycle.
- Lottery scheduling has no global state so incorporating new processes is a breeze
 - In stride scheduling, its difficult to assign a correct pass value to the new job
 - If the new job is given a pass value of 0, it monopolizes the CPU

Completely Fair Scheduler (CFS)

- Linux Scheduler
- Uses Fair Share Scheduling
- Doesn't have a fixed time slice
- Keeps a counter vruntime for each process and adds a time slice to it for each process
- The value of time slice is dynamically adjusted for the next run based on the available workload
- When a scheduling decision occurs, it'd select the process with lowest vruntime to run
- Run for how long?
 - Time slice too small → good fairness but scheduling overhead
 - Time slice too big → low scheduling overhead but costs fairness in short term

CFS

- Maintains various parameters
 - sched_latency
 - $time_slice = \frac{sched_latency}{num_processes}$
 - sched_latency = 48, num_processes=4
 - time_slice = 12




Figure 9.4: CFS Simple Example

CFS

- What if too many processes running?
 - min_granularity
 - $\text{time_slice} = \max(\text{min_granularity}, \text{calculated_time_slice_value})$
- Timer interrupt goes off every 1 ms
 - Everytime a process leaves CPU, its vruntime gets updated
 - This way eventually fairness is ensured

CFS: Niceness

- Processes can be nice to other processes (or not)
- A higher nice value would give them a lower weight \rightarrow smaller slice
- `nice_A = -5`, `nice_B = 0` (default value)

$$\text{time_slice}_k = \frac{\text{weight}_k}{\sum_{i=0}^{n-1} \text{weight}_i} \cdot \text{sched_latency} \quad (9.1)$$

$$\text{vruntime}_i = \text{vruntime}_i + \frac{\text{weight}_0}{\text{weight}_i} \cdot \text{runtime}_i \quad (9.2)$$

- process vruntime will also increase in inverse proportion to its weight

Data structure: RB Tree

- Self balancing and Sorted
 - Removing a process is trivial
 - Insertion is $\log(n)$
- Only has ready processes inserted on their vruntime values
- I/O, sleeping processes assigned min vruntime when they join.




Figure 9.5: CFS Red-Black Tree

Summary

- We introduced proportional-share scheduling and briefly discussed three approaches
 - lottery scheduling, stride scheduling, and the Completely Fair Scheduler (CFS)
- Lottery uses randomness to achieve proportional share, is non-deterministic but scales well
- Stride uses stride to calculate exact proportions hence is deterministic but does not scale well
- CFS is like a weighted round robin scheduler with dynamic time slices, scales and performs well