

Operating System (OS)

CS232

Memory Virtualization: Address Translation

Dr. Muhammad Mobeen Movania

Outlines

- Limited Direct Execution (LDE)
- Address Translation
- Memory Virtualization Assumptions
- Memory Accesses and Virtualization
- Dynamic Relocation and Address Translation Examples
- Hardware Support and OS issues
- Dynamic Relocation Issues
- Summary

Limited Direct Execution (LDE)

- OS allows the program to run directly on the hardware
 - OS intervenes only when
 - something bad happens like an exception or
 - the process's time slice expires raising a timer interrupt
- Why LDE?
 - Efficiency (sharing of memory and resources among processes)
 - Control (restrict process to its own memory)

Address Translation

- OS has to ensure that the virtualization of memory through address translation must provide
 - Efficiency (using hardware support)
 - Control (for isolation of process memory)
- The hardware transforms each memory access (e.g., an instruction fetch, load, or store), changing **virtual** address to a **physical** address
- OS must keep information of used and free space for allocating processes

Memory Virtualization Assumptions

- A process address space lies contiguously in memory
 - Helps in accessing data quicker
- Size of a process address space is smaller than physical memory
 - Simpler to deal with
- All process address spaces are of same size
 - Access to a process is easier and more **direct**

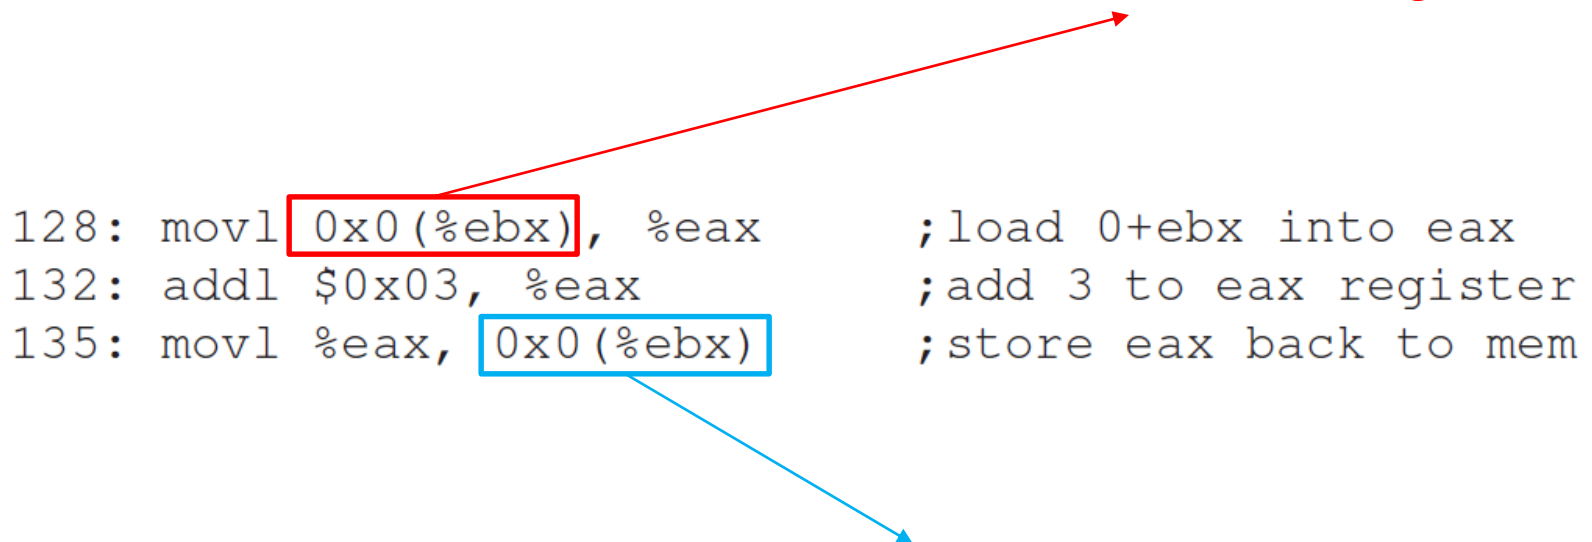
Example

```
void func() {  
    int x = 3000;  
    x = x + 3;  
    ...  
}
```

ebx register contains 15KB which is
(address of x in stack memory)

Read from address stored in ebx register

```
128: movl 0x0(%ebx), %eax    ;load 0+ebx into eax  
132: addl $0x03, %eax        ;add 3 to eax register  
135: movl %eax, 0x0(%ebx)    ;store eax back to mem
```



Write to address stored in ebx register


Memory Accesses

- Fetch instruction at address 128
- Execute this instruction (load from address 15 KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)



Virtualization

- From the process perspective, its address space starts at 0
- In reality the OS might not place the process image at that address
- How can the OS provide this illusion in a manner transparent to the process?



Dynamic relocation

- Two special registers in CPU
 - Base
 - Bound (limit)
- Every address generated by the program is translated by the hardware:

$$\text{physical address} = \text{virtual address} + \text{base}$$

- **Memory Management Unit (MMU)**: part of CPU that helps with address translation

Memory Management Unit (MMU)

- Base register
 - Contains the start of process address space in physical memory
 - Used in address translation
- Bound register
 - Can contain either:
 - The size of process address space, or
 - The final physical address of the process address space
 - Used in checking for illegal memory accesses

Address Translation Example

- Process with an address space of 4KB
 - Min Virtual Address: 0 bytes
 - Max Virtual Address: $4 * 1024 = 4096$ bytes

Virtual address	Physical Address
0	16KB= 16384 bytes
1KB	17KB
3000 bytes	16KB+3000=19384 bytes
4400 bytes	16KB+4400=20784 bytes Fault(Out of bounds memory access)

H/W support we need for MM

Hardware Requirements	Notes
Privileged mode	<i>Needed to prevent user-mode processes from executing privileged operations</i>
Base/bounds registers	<i>Need pair of registers per CPU to support address translation and bounds checks</i>
Ability to translate virtual addresses and check if within bounds	<i>Circuitry to do translations and check limits; in this case, quite simple</i>
Privileged instruction(s) to update base/bounds	<i>OS must be able to set these values before letting a user program run</i>
Privileged instruction(s) to register exception handlers	<i>OS must be able to tell hardware what code to run if exception occurs</i>
Ability to raise exceptions	<i>When processes try to access privileged instructions or out-of-bounds memory</i>

OS support we need for MM

OS Requirements	Notes
Memory management	<i>Need to allocate memory for new processes; Reclaim memory from terminated processes; Generally manage memory via free list</i>
Base/bounds management	<i>Must set base/bounds properly upon context switch</i>
Exception handling	<i>Code to run when exceptions arise; likely action is to terminate offending process</i>


Can the OS move a process image in physical memory?

Operating System Responsibilities

- OS must
 - *keep track of free space* in a data structure called **free list**.
 - *do termination housekeeping* when a process is terminated (reclaim process's memory)
 - *save and restore* the base-and-bounds pair in **PCB** when it switches between processes.
 - provide **exception handlers** to handle exceptions
- OS just sets up the hardware and lets the process run directly on the CPU
 - only when the process **misbehaves** does the OS have to become involved.

Dynamic Relocation Issues

- Internal fragmentation



Summary

- We have extended the concept of limited direct execution with **address translation**
- OS can control each and every memory access ensuring the accesses stay within the bounds of the address space
- Efficiency depends on hardware support
- Address translation remains transparent to a process