

Operating System (OS)

CS232

Memory Management: Smaller Page Tables

Dr. Muhammad Mobeen Movania

Outlines

- Issues with Page Tables
- Bigger Pages
- Hybrid Approach
- Multi-level Page Tables
- Example, Advantages and Disadvantages
- Inverted Page Tables
- Summary

Issues with Page Tables

- Page tables are linear structures that rapidly consume a lot of memory
- Example
 - 32-bit address space ($2^{32}=4294967296$ bytes), $4\text{KB}=2^{12}$ page size
 - 4 bytes page table entry
 - Total pages = $2^{32}/2^{12}=2^{20}=1048576$
 - Total page table size = $2^{20}*4=4194304=4\text{MB}$
 - We have one page table per process, if hundred processes running in system, memory requirements will increase drastically
- Major Concern
 - How can we reduce the page table size?

Simple Solution: Bigger Pages

- Make pages bigger, less page table entries and less storage requirement of page tables
- Example
 - 32-bit address space ($2^{32}=4294967296$ bytes),
16KB= 2^{14} page size
 - 4 bytes page table entry
 - Total pages= $2^{32}/2^{14}=2^{18}=262144$
 - Total page table size = $2^{18}*4=1048576=1\text{MB}$
 - Reduce page table size to a quarter as compared to that with page of 4 KB size

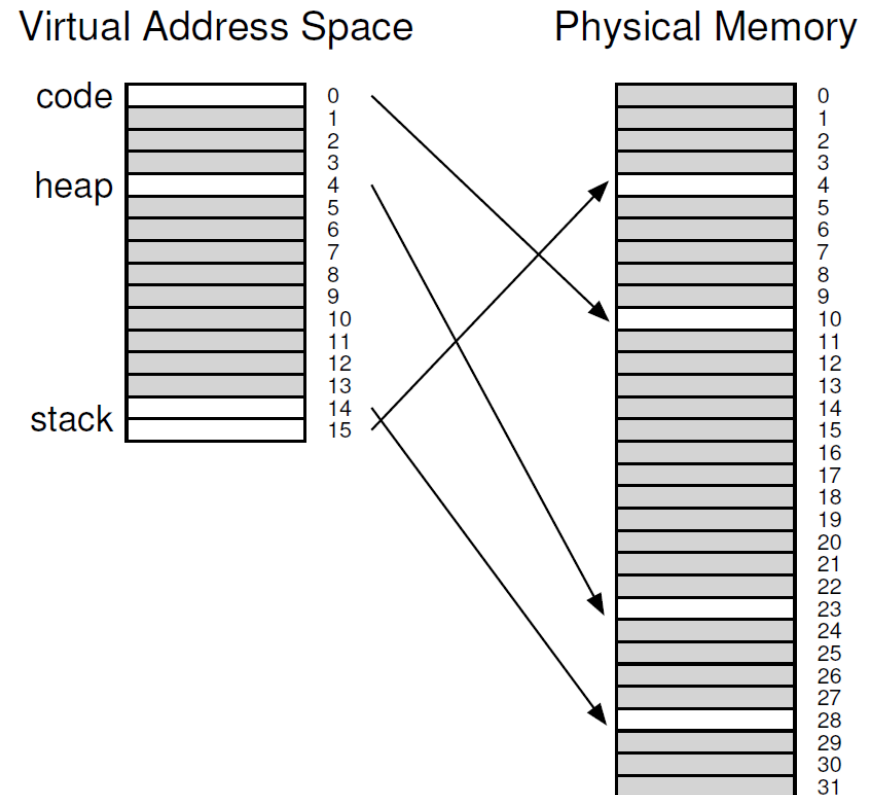
Second Solution: Hybrid Approach

- Key Observation
 - Most of the space inside page table is free
 - Combine segmentation with paging to only store entries with valid data
- Key Idea
 - Have one page table per logical segment
 - Base register: contains physical address of page table of the segment
 - Bounds register: indicates the end of the page table

Example 1


- 16KB address space with 1KB pages and small segments

PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
23	1	rw-	1	1
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
28	1	rw-	1	1
4	1	rw-	1	1



Example 2

- 32-bit virtual address space split into three segments (need 3 base/bounds pair) and 4KB pages



00	Unused
01	Code segment
10	Heap segment
11	Stack segment

```
SN          = (VirtualAddress & SEG_MASK) >> SN_SHIFT
VPN         = (VirtualAddress & VPN_MASK) >> VPN_SHIFT
AddressOfPTE = Base[SN] + (VPN * sizeof(PTE))
```

Hybrid Approach – Pros and Cons

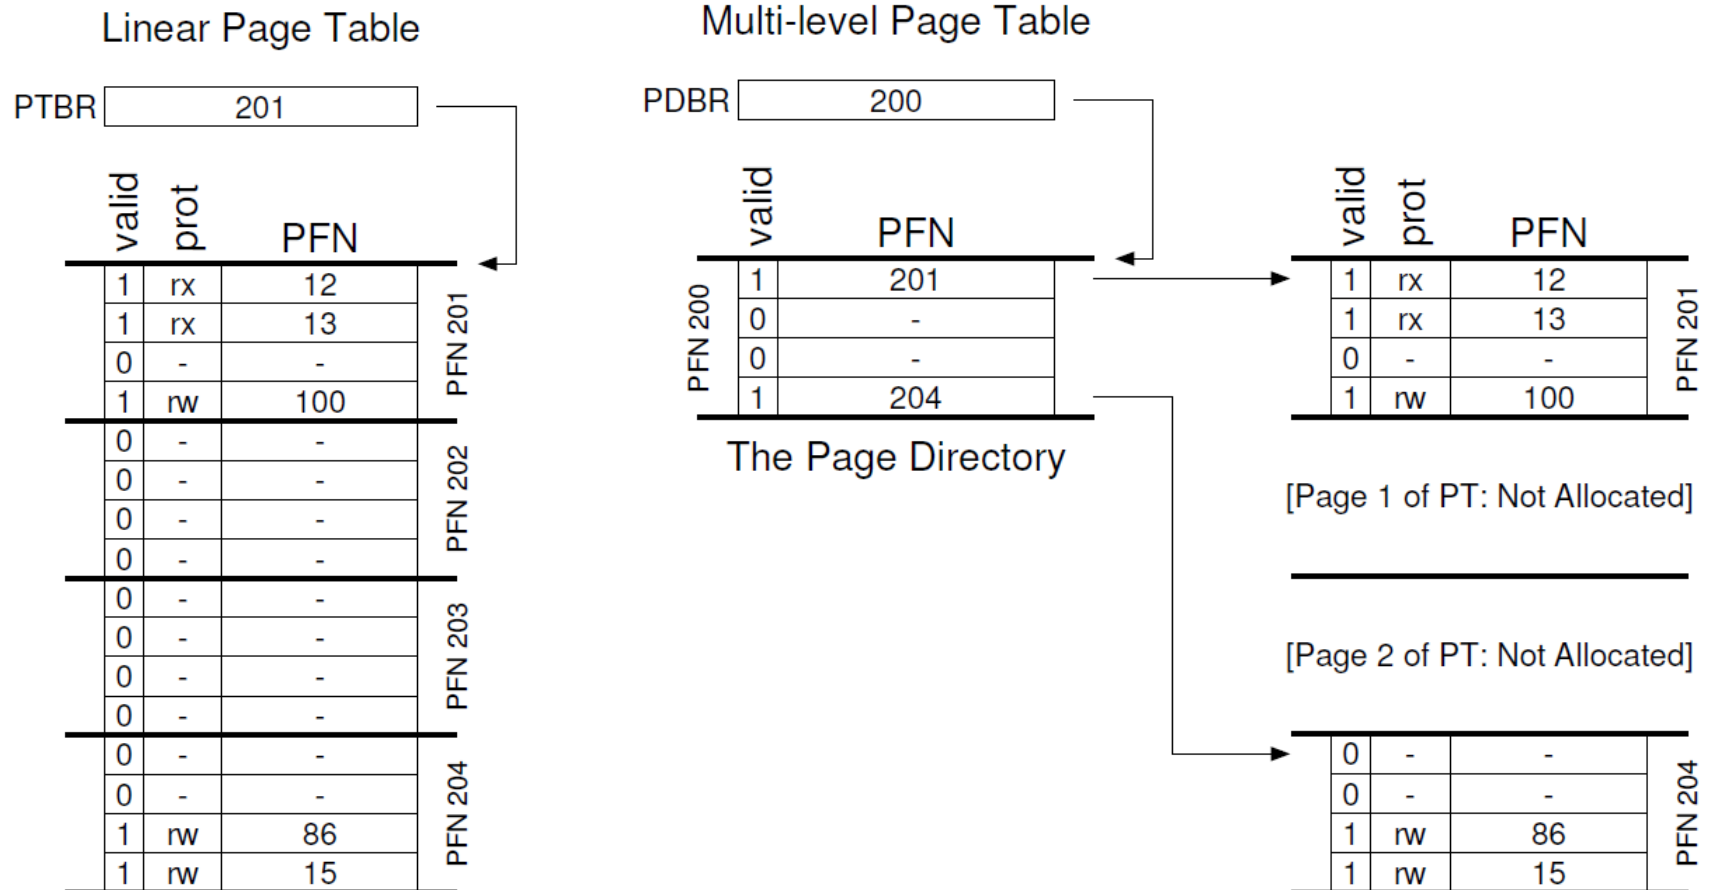
- Pros
 - Reduces page table sizes
- Cons
 - Have multiple base and bound registers
 - Dividing address space into fixed size portions is not flexible
 - Variable sized page tables depending on the segment size

Third Solution – Multi-level page tables

- Chop up the *page table* into page sized fragments
- If an entire page of page-table entries is invalid, don't allocate space for that page in page table RAM
- A new structure, page directory, tells you:
 - which *pages of the page table* are valid!
 - where in RAM would you find the valid pages of the page table

Example – Multi-level Page Tables

- Comparing linear PT (L) vs multi-level PT (R)



Multi-level Page Tables – Pros and Cons

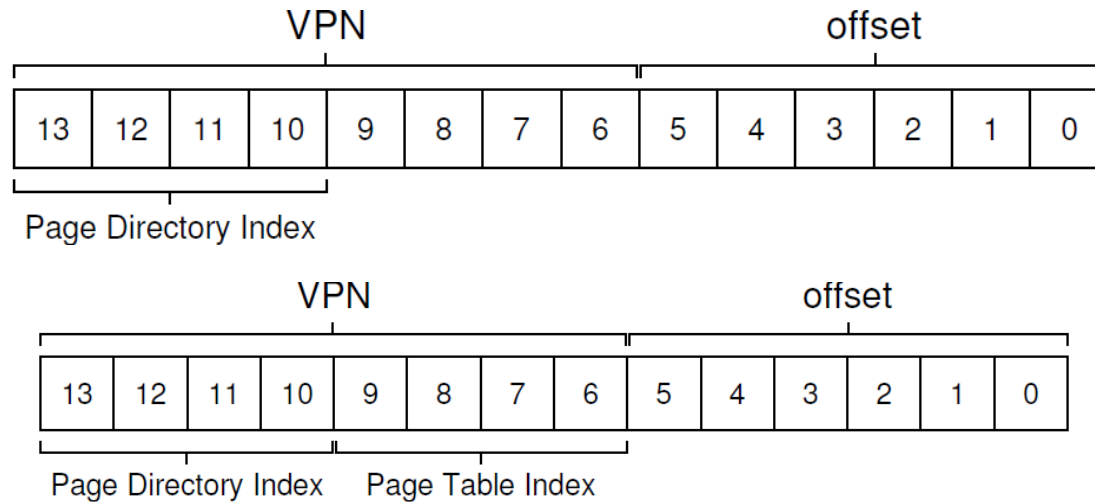
- Pros
 - Page table sizes are reduced
 - Page tables can be distributed in non contiguous pages
 - Previous schemes required contiguous series of pages in memory
- Cons
 - Complexity of implementation
 - On a TLB miss, we've to access the RAM 3 times
 - example of Time-space tradeoff

Example: Multi-level Page Table

- 16KB address space, 64 byte pages
- Linear page table size
 - 16 KB address space ($2^{14}=16384$ bytes)
 - 64 byte page size = 2^6
 - Total pages= $2^{14}/2^6=2^8=256$
 - Total page table size = $2^8*4=1024=1\text{KB}$
 - Page table split into 16, 64 byte pages (16 PTEs)

0000 0000	code
0000 0001	code
0000 0010	(free)
0000 0011	(free)
0000 0100	heap
0000 0101	heap
0000 0110	(free)
0000 0111	(free)
.....	... all free ...
1111 1100	(free)
1111 1101	(free)
1111 1110	stack
1111 1111	stack

Example: Multi-level Page Table



PDEAddr = PDBaseAddr + (PDIndex * sizeof(PDE))

PTEAdd = (PDE.PFN<<SHIFT) + (PTIndex*sizeof(PTE))

Example: Multi-level Page Table

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

Example: Multi-level Page Table

- Crux
 - Instead of allocating 16 pages of a linear page table, we allocate 3 pages
 - One for page directory
 - Two for page table chunks (PFN:100, PFN:101)


Multi-level Page Table- Algorithm

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset    = VirtualAddress & OFFSET_MASK
6          PhysAddr  = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else                    // TLB Miss
11     // first, get page directory entry
12     PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13     PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14     PDE      = AccessMemory(PDEAddr)
15     if (PDE.Valid == False)
16         RaiseException(SEGMENTATION_FAULT)
17     else
18         // PDE is valid: now fetch PTE from page table
19         PTIndex = (VPN & PT_MASK) >> PT_SHIFT
20         PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))
21         PTE      = AccessMemory(PTEAddr)
22         if (PTE.Valid == False)
23             RaiseException(SEGMENTATION_FAULT)
24         else if (CanAccess(PTE.ProtectBits) == False)
25             RaiseException(PROTECTION_FAULT)
26         else
27             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
28             RetryInstruction()
```


Figure 20.6: Multi-level Page Table Control Flow

Multi-level page tables – more than 2 levels

- 30-bit address space, 512 byte pages



- Page directory doesn't fit in one page



Inverted Page Tables

- One page table for the whole system
- Indices are PFNs instead of VPNs
- PTE contains: PID, VPN
- Have to do a linear search
 - Hash tables may help reduce search time

Example: Inverted Page Table

[View Shockwave Animation](#)

Summary

- We saw how non-linear page tables are built
- Small tables
 - used in memory constrained systems
- Large tables
 - used with a reasonable amount of memory and with workloads that actively use a large number of pages help speeds up TLB misses
- Leftover issues
 - What if size of total number of pages in a system exceeds RAM?
 - What if there are so many page tables that the RAM space allocated to kernel falls short?