

Operating System (OS)

CS232

Beyond Physical Memory: Mechanisms

Dr. Muhammad Mobeen Movania

Outlines

- Illusion of large virtual address space
- What is swap space?
- The Present Bit
- What happens at page fault?
- What if memory is full?
- Page fault control flow
- When replacements really occur?
- Summary

Illusion of large virtual address space

- Usually, we assume that process address space can fit within the available physical memory
- To give illusion of large virtual address space, OS uses the next level of memory hierarchy, the hard disk
 - Pages are moved to hard disk if they can't be accommodated in system memory
 - Slow system response if there is a lot of page swapping

Swap Space

- Some space on the hard disk for moving pages back and forth from main memory

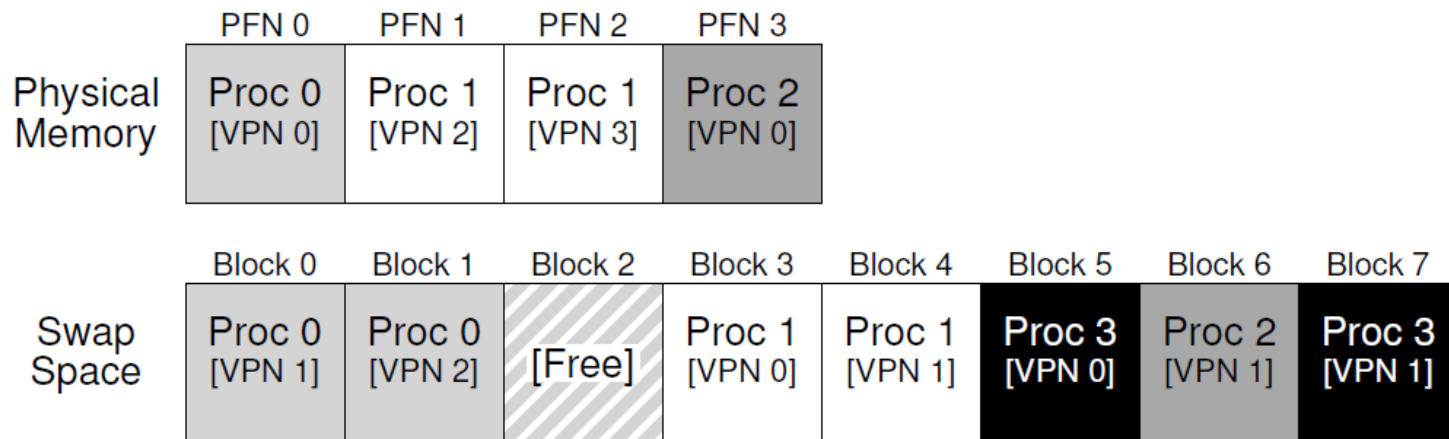


Figure 21.1: Physical Memory and Swap Space

The Present Bit

- A bit in the page table entry
- Tells if this page is in physical memory or on disk




Figure 18.5: An x86 Page Table Entry (PTE)

- Accessing a page which is not in physical memory is called ***page fault***

What happens at page fault?

- Irrespective of hardware or software managed TLB, OS handles page fault
 - OS page fault handler runs
 - If a page is not present and has been swapped to disk, the OS will need to swap the page into memory in order to service the page fault
 - When the disk I/O completes, the OS will then
 - update the page table to mark the page as present
 - update the PFN field of the page-table entry (PTE) to record the in-memory location of the newly-fetched page
 - Finally, retry the instruction

What if memory is full?

- OS might like to first **page out** one or more pages to make room for the new page(s) the OS is about to bring in.
- The process of picking a page to kick out, or **replace** is known as the **page-replacement policy**.

Page Fault Control Flow

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else    // TLB Miss
11     PTEAddr = PTBR + (VPN * sizeof(PTE))
12     PTE = AccessMemory(PTEAddr)
13     if (PTE.Valid == False)
14         RaiseException(SEGMENTATION_FAULT)
15     else
16         if (CanAccess(PTE.ProtectBits) == False)
17             RaiseException(PROTECTION_FAULT)
18         else if (PTE.Present == True)
19             // assuming hardware-managed TLB
20             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
21             RetryInstruction()
22         else if (PTE.Present == False)
23             RaiseException(PAGE_FAULT)
```

Figure 21.2: Page-Fault Control Flow Algorithm (Hardware)

When replacements really occur?

- Most operating systems thus have some kind of **high watermark** (HW) and **low watermark** (LW) to help decide when to start evicting pages from memory
 - when the OS notices that there are fewer than LW pages available, a background thread (swap daemon) is responsible for freeing memory runs
 - It evicts pages until there are HW pages available

Summary

- We have introduced the notion of accessing more memory than is physically present within a system
- Page table must have a bit (**present bit**) to tell if a requested page is in memory or not
- If the page is not present
 - OS **page-fault handler** runs to service the **page fault**, and thus arranges for the transfer of the desired page from disk to memory
- All of this happens transparent to the process