

Operating System (OS)

CS232

Concurrency: Mutual Exclusion, Locks

Dr. Muhammad Mobeen Movania

Dr Muhammad Saeed

Outlines

- Concurrency & Race Condition (recap)
- Producer Consumer Problem
 - Mutual exclusion
 - Synchronization
- Semaphore
 - Counting and binary
- Mutual Exclusion
 - Mutex
 - Race Condition – solution
- Lock based Concurrent Data Structures

Concurrency & Race Condition

```
long *balance; //Shared Memory among multiple processes
```

```
void credit(int arg){  
    int amount = arg;  
    printf("Credit : balance = balance + %d\n", arg);  
    for(long i=0; i<5000000; i++){  
        *balance = *balance + amount;  
    }  
}
```

```
void debit(int arg){  
    int amount = arg;  
    printf("Debit : balance = balance - %d\n", arg);  
    for(long i=0; i<5000000; i++){  
        *balance = *balance - amount;  
    }  
}
```

```
int main(){  
    //IPC - Shared Memory technique is used to store variable balance  
    key_t key = ftok("sm_bal", 65);  
    int shm_id=shmget(key, 8, IPC_CREAT | 0666);  
    balance = (long*)shmat(shm_id, NULL, 0);  
    *balance=0; //initializing balance
```

```
    int cpid = fork();  
    if (cpid == 0){  
        credit(1);  
        shmdt(balance);  
        exit(0);  
    }
```

```
    else{  
        debit(1);  
        waitpid(cpid, NULL, 0);  
        //debit(1);  
        printf("Value of balance is: %ld\n", *balance);  
        shmdt(balance);  
        shmctl(shm_id, IPC_RMID, NULL);  
        return 0;  
    }
```

```
Debit : balance = balance - 1  
Credit : balance = balance + 1  
Value of balance is: -3762270
```

Concurrency & Race Condition

```
long balance = 0; // Shared, Global Variable
```

```
void * credit(void * arg){
    int amount = (int) arg;
    printf("Credit : balance = balance + %d\n", (int) arg);
    for(long i=0; i<5000000; i++){
        balance = balance + amount;
    }
    pthread_exit(NULL);
}
```

```
void * debit(void * arg){
    int amount = (int) arg;
    printf("Debit : balance = balance - %d\n", (int) arg);
    for(long i=0; i<5000000; i++){
        balance = balance - amount;
    }
    pthread_exit(NULL);
}
```

```
int main(){
    pthread_t t1, t2;
    pthread_create(&t1, NULL, credit, (void *)5);
    pthread_create(&t2, NULL, debit, (void *)5);

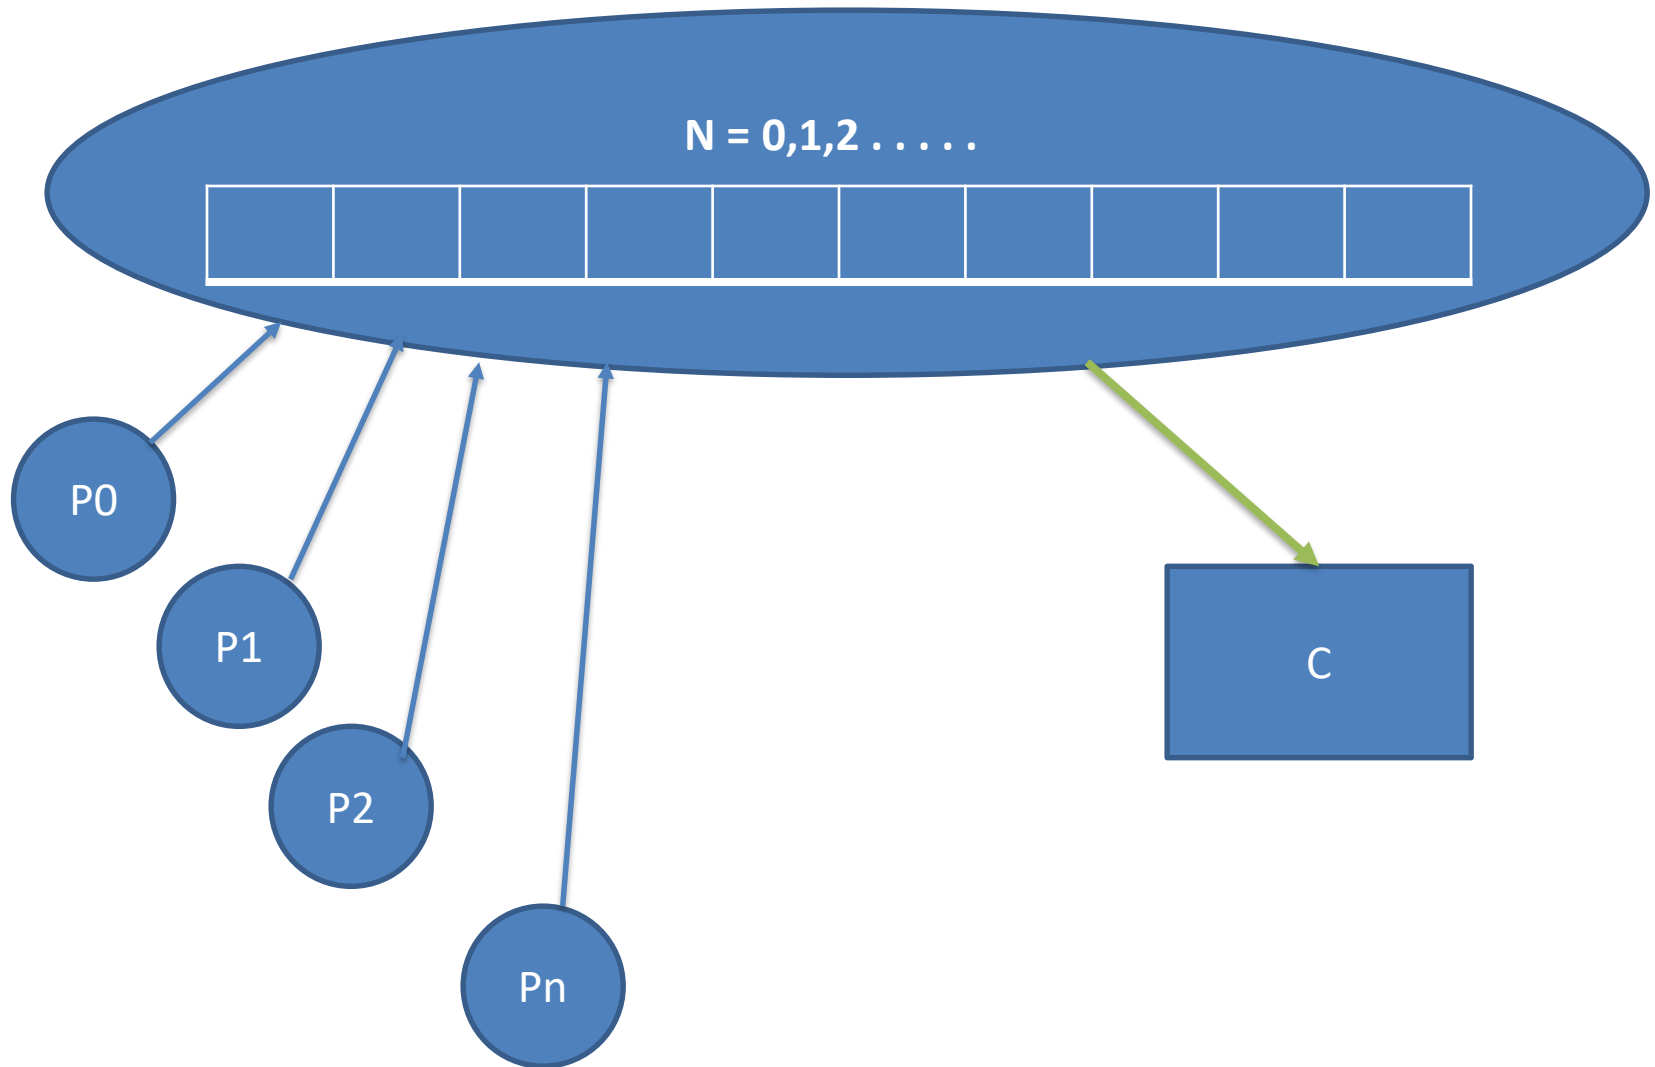
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Value of balance is : %ld\n", balance);
    return 0;
}
```

```
Debit : balance = balance - 5
Credit : balance = balance + 5
Value of balance is :-417255
```

Larger Context

- Client/Server Applications
- Peer to Peer Applications
- Producer / Consumer Problem
- Readers / Writers Problem

Produce/ Consumer Problem



Solution

- Requirements
 - Mutual Exclusion
 - Synchronization
- Semaphore
- Programming Techniques
 - Mutex & Condition Variables
 - Semaphore

Semaphore (counting)

```
struct semaphore {  
    int count;  
    queueType queue;  
};  
  
void semWait(semaphore s)  
{  
    s.count--;  
    if (s.count < 0) {  
        /* place this process in s.queue */;  
        /* block this process */;  
    }  
}  
  
void semSignal(semaphore s)  
{  
    s.count++;  
    if (s.count <= 0) {  
        /* remove a process P from s.queue */;  
        /* place process P on ready list */;  
    }  
}
```


Semaphore (binary)

```
struct binary_semaphore {  
    enum {zero, one} value;  
    queueType queue;  
};
```

```
void semWaitB(binary_semaphore s)  
{  
    if (s.value == one)  
        s.value = zero;  
    else {  
        /* place this process in s.queue  
        /* block this process */;  
    }  
}
```

```
void semSignalB(semaphore s)  
{  
    if (s.queue is empty())  
        s.value = one;  
    else {  
        /* remove a process P from s.queue  
        /* place process P on ready list */;  
    }  
}
```

Producer / Consumer - Solution

```
/* program producerconsumer */
int n;
binary_semaphore s = 1, delay = 0;
void producer()
{
    while (true) {
        produce();
        semWaitB(s);
        append();
        n++;
        if (n==1) semSignalB(delay);
        semSignalB(s);
    }
}
void consumer()
{
    int m; /* a local variable */
    semWaitB(delay);
    while (true) {
        semWaitB(s);
        take();
        n--;
        m = n;
        semSignalB(s);
        consume();
        if (m==0) semWaitB(delay);
    }
}
void main()
{
    n = 0;
    parbegin (producer, consumer);
}
```

Producer / Consumer - Solution

```
/* program producerconsumer */
semaphore n = 0, s = 1;
void producer()
{
    while (true) {
        produce();
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}
void consumer()
{
    while (true) {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        consume();
    }
}
void main()
{
    parbegin (producer, consumer);
}
```

Figure 5.11 A Solution to the Infinite-Buffer Producer/Consumer Problem
Using Semaphores

Mutual Exclusion

- To protect shared resources from race condition and data inconsistency. (Balance = Balance +/- Amount)

Thread 1	Thread 2	Shared Data X
A = X		100
	B = X	100
	B = B + 300	100
A = A + 200		100
X = A		300
	X = B	400

Thread - Issues

- When we have a number of threads accessing a shared resource
 - Multiple threads can enter critical section simultaneously
 - Critical section may be interrupted in the middle
- How can we make the critical section mutually exclusive!
 - i.e. only one thread can enter a critical section at any given instant
 - This would have the effect of the critical section being atomic!
- Solution
 - Locks

Mutex

- Library - #include <pthread.h>
- Structure
 - pthread_mutex_t
- Functions
 - pthread_mutex_init (mutex , attr)
 - pthread_mutex_destroy (mutex)
 - pthread_mutexattr_init (attr)
 - pthread_mutexattr_destroy (attr)

Locking and Unlocking Mutex

- Functions
 - `pthread_mutex_lock (mutex)`
 - `pthread_mutex_unlock (mutex)`
 - `pthread_mutex_trylock (mutex)`

Race Condition - Solution

```
//For MUTEX LOCK()/UNLOCK()
pthread_mutex_t mut = PTHREAD_MUTEX_INITIALIZER;
//Shared Variable
long balance = 0; //Shared, Global Variable
```

```
void * credit(void * arg){
    int amount = (int) arg;
    printf("Credit : balance = balance + %d\n", (int) arg);
    for(long i=0; i<5000000; i++){
        pthread_mutex_lock(&mut);
        balance = balance + amount;
        pthread_mutex_unlock(&mut);
    }
    pthread_exit(NULL);
}
```

```
void * debit(void * arg){
    int amount = (int) arg;
    printf("Dedit : balance = balance - %d\n", (int) arg);
    for(long i=0; i<5000000; i++){
        pthread_mutex_lock(&mut);
        balance = balance - amount;
        pthread_mutex_unlock(&mut);
    }
    pthread_exit(NULL);
}
```

```
int main(){
    pthread_t t1, t2;
    pthread_create(&t1, NULL, credit, (void *)5);
    pthread_create(&t2, NULL, debit, (void *)5);

    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Value of balance is : %ld\n", balance);
    return 0;
}
```

```
Dedit : balance = balance - 5
Credit : balance = balance + 5
Value of balance is : 0
```


Reading Assignment

- Types of locks and locking strategies
- Their pros and cons
- Two phase and Three phase locking

Locks

- Locks are constructs provided by the OS (or libraries)
- Locks can be acquired and released
- With OS (and hardware) support, it is ensured that only one thread can acquire a given lock at any given time!
- We use locks to protect our critical sections in multiple threads

```
1  lock_t mutex; // some globally-allocated lock 'mutex'
2  ...
3  lock(&mutex);
4  balance = balance + 1;
5  unlock(&mutex);
```

Locks (2)

- A lock is a variable which can be in any of two states:
 - Available (or free, or unlocked)
 - Acquired (or held, or locked)
- When a thread calls `lock()` on a particular lock variable
 - if the lock is in free state, it acquires the lock, and `lock()` function returns
 - If the lock is in held state, it will block and the `lock()` function will not return until it has acquired the lock
- A thread can free an acquired lock by calling `unlock()` on it.

Lock (3)

- Two approaches to locking
 - **Coarse-grained locking:** Having one big lock used when any critical section is accessed
 - **Fine-grained locking:** Having different data and data structures with different locks
- Pthread locks (mutexes)

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3  Pthread_mutex_lock(&lock); // wrapper; exits on failure  
4  balance = balance + 1;  
5  Pthread_mutex_unlock(&lock);
```

Evaluating Locks

- Goals:
 - Correctness - Achieve mutual exclusion
 - Fairness – Give chance to each waiting thread
 - Performance - Be efficient
 - Single thread single processor
 - Multiple threads single processor
 - Multiple threads multiple processors

Lock Implementation-Disable Interrupt

- First solution
 - Disable interrupts
 - CPU has special instructions for this
 - Works for single processor systems
- Issues
 - Turning off interrupts is a privileged operation
 - If interrupts are off, useful interrupts can be lost

```
1  void lock() {  
2      DisableInterrupts();  
3  }  
4  void unlock() {  
5      EnableInterrupts();  
6  }
```

Lock Implementation – Use Flag

- Second solution
 - Use a variable (flag) to communicate b/w threads

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 -> lock is available, 1 -> held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1;          // now SET it!
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Lock Implementation – Use Flag

- Second Solution
 - Issues
 - Performance (spin waiting)
 - Correctness? No Mutual Exclusion

Thread 1

```
call lock ()  
while (flag == 1)  
interrupt: switch to Thread 2
```

```
flag = 1; // set flag to 1 (too!)
```

Thread 2

```
call lock ()  
while (flag == 1)  
flag = 1;  
interrupt: switch to Thread 1
```

Figure 28.2: Trace: No Mutual Exclusion

Lock Implementation - (Use H/W Support)

- Accessing a shared flag might be interrupted hence mutual exclusion may not be possible
- Third solution – Use hardware support
 - test-and-set (atomic exchange) instruction
 - It provide an atomic instruction

```
1  int TestAndSet(int *old_ptr, int new) {  
2      int old = *old_ptr; // fetch old value at old_ptr  
3      *old_ptr = new;      // store 'new' into old_ptr  
4      return old;          // return the old value  
5  }
```

Lock Implementation - (Use H/W Support)

- Test-and-set spin lock

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0: lock is available, 1: lock is held
7      lock->flag = 0;
8  }
9
10 void lock(lock_t *lock) {
11     while (TestAndSet(&lock->flag, 1) == 1)
12         ; // spin-wait (do nothing)
13 }
14
15 void unlock(lock_t *lock) {
16     lock->flag = 0;
17 }
```

- Issues

- Needs a preemptive scheduler on a single processor otherwise a thread may never relinquish the CPU

Evaluating spin locks

- Correctness
 - Yes it provides mutual exclusion
- Fairness
 - No, spin locks don't provide fairness guarantee (a thread spinning may lock the CPU causing waiting threads to starve)
- Performance
 - Worse performance if the thread holding the lock is preempted within a critical section
 - All other threads will wait spin wasting CPU cycles

Lock Implementation – (Use H/W Support)

- compare-and-swap instruction

```
1  int CompareAndSwap(int *ptr, int expected, int new) {  
2      int actual = *ptr;  
3      if (actual == expected)  
4          *ptr = new;  
5      return actual;  
6  }
```

```
1  void lock(lock_t *lock) {  
2      while (CompareAndSwap(&lock->flag, 0, 1) == 1)  
3          ; // spin  
4  }
```

Lock Implementation – (Use H/W Support)

- load-linked and store-conditional instruction
- Both instructions operate in tandem

```
1  int LoadLinked(int *ptr) {
2      return *ptr;
3  }
4
5  int StoreConditional(int *ptr, int value) {
6      if (no update to *ptr since LoadLinked to this address) {
7          *ptr = value;
8          return 1; // success!
9      } else {
10         return 0; // failed to update
11     }
12 }
```

Lock Implementation – (Use H/W Support)

- Store-conditional succeeds only if no intervening store has happened to that address since the last load-linked!!

```
1 void lock(lock_t *lock) {
2     while (1) {
3         while (LoadLinked(&lock->flag) == 1)
4             ; // spin until it's zero
5         if (StoreConditional(&lock->flag, 1) == 1)
6             return; // if set-it-to-1 was a success: all done
7                     // otherwise: try it all over again
8     }
9 }
10
11 void unlock(lock_t *lock) {
12     lock->flag = 0;
13 }
```

Lock Implementation – (Use H/W Support)

- fetch-and-add instruction
- Used to implement ticket lock

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void lock_init(lock_t *lock) {
7      lock->ticket = 0;
8      lock->turn   = 0;
9  }
10
11 void lock(lock_t *lock) {
12     int myturn = FetchAndAdd(&lock->ticket);
13     while (lock->turn != myturn)
14         ; // spin
15 }
16
17 void unlock(lock_t *lock) {
18     lock->turn = lock->turn + 1;
19 }
```

```
1  int FetchAndAdd(int *ptr) {
2      int old = *ptr;
3      *ptr = old + 1;
4      return old;
5  }
```

Figure 28.7: Ticket Locks

Lock Implementation – (Use H/W Support)

- All lock implementations using hardware support spin wait wasting CPU cycles
- We can avoid spin wait by process calling **yield** function to voluntarily give the CPU time to other waiting thread
- Requires support of the operating system, yield is a system call

Lock Implementation – Using yield

```
1  void init() {
2      flag = 0;
3  }
4
5  void lock() {
6      while (TestAndSet(&flag, 1) == 1)
7          yield(); // give up the CPU
8  }
9
10 void unlock() {
11     flag = 0;
12 }
```

Locks Implementation – Use Queues

- Use Queues

```
13 void lock(lock_t *m) {
14     while (TestAndSet(&m->guard, 1) == 1)
15         ; //acquire guard lock by spinning
16     if (m->flag == 0) {
17         m->flag = 1; // lock is acquired
18         m->guard = 0;
19     } else {
20         queue_add(m->q, gettid());
21         m->guard = 0;
22         park();
23     }
24 }
```

```
26 void unlock(lock_t *m) {
27     while (TestAndSet(&m->guard, 1) == 1)
28         ; //acquire guard lock by spinning
29     if (queue_empty(m->q))
30         m->flag = 0; // let go of lock; no one wants it
31     else
32         unpark(queue_remove(m->q)); // hold lock
33                                     // (for next thread!)
34     m->guard = 0;
35 }
```

```
1  typedef struct __lock_t {
2      int flag;
3      int guard;
4      queue_t *q;
5  } lock_t;
6
7  void lock_init(lock_t *m) {
8      m->flag = 0;
9      m->guard = 0;
10     queue_init(m->q);
11 }
```

- Sleep instead of spinning, if lock is held.
- Park() and unpark() support provided by Solaris for sleep
- *guard* is a spin lock around *flag* and *wait queues*
- some spinning is done but only for the time we access the flag & queue
- Flag is not set to 0 on unpark()!

Futexes

```
1 void mutex_lock (int *mutex) {
2     int v;
3     /* Bit 31 was clear, we got the mutex (the fastpath) */
4     if (atomic_bit_test_set (mutex, 31) == 0)
5         return;
6     atomic_increment (mutex);
7     while (1) {
8         if (atomic_bit_test_set (mutex, 31) == 0) {
9             atomic_decrement (mutex);
0             return;
1         }
2         /* We have to waitFirst make sure the futex value
3            we are monitoring is truly negative (locked). */
4         v = *mutex;
5         if (v >= 0)
6             continue;
7         futex_wait (mutex, v);
8     }
9 }
```

Futexes

- MSB is the status
- All other bits keep count

```
21 void mutex_unlock (int *mutex) {  
22     /* Adding 0x80000000 to counter results in 0 if and  
23        only if there are not other interested threads */  
24     if (atomic_add_zero (mutex, 0x80000000))  
25         return;  
26  
27     /* There are other threads waiting for this mutex,  
28        wake one of them up.  */  
29     futex_wake (mutex);  
30 }
```

Two-phased locks

- If the lock is held, spin for a while to see if it's acquirable in near future.
- If not, then go to sleep.

Summary

- We saw what are the different mechanisms of locks and how are they implemented
- We saw the role that hardware support provides in lock implementation
- We saw the different metrics on which a lock implementation is evaluated (correctness, fairness and performance)
- We also saw how spin locks waste CPU cycles and so alternate OS supported functions like yield and sleep are used to improved performance