

# Operating System (OS)

## CS232

Concurrency: Lock-based Concurrent Data Structures

Dr. Muhammad Mobeen Movania  
Dr Muhammad Saeed

# Outlines

- What is thread safety of a data structure?
- Concurrent Counter
- Issues and resolution of Concurrent Counter
- Concurrent Linked Lists
- Concurrent Queues
- Concurrent Hash Tables
- Summary

# What is thread safety?

- Adding locks to a data structure to make it usable by multiple threads concurrently makes the data structure **thread safe**
- How locks are added determines correctness and performance of the data structure
- Many different data structures exist, simplest is a concurrent counter

# Concurrent Counter

- Simple counter without locks is not thread safe

```
1  typedef struct __counter_t {  
2      int value;  
3  } counter_t;  
4  
5  void init(counter_t *c) {  
6      c->value = 0;  
7  }  
8  
9  void increment(counter_t *c) {  
10     c->value++;  
11 }  
12  
13 void decrement(counter_t *c) {  
14     c->value--;  
15 }  
16  
17 int get(counter_t *c) {  
18     return c->value;  
19 }
```

Figure 29.1: A Counter Without Locks

# Concurrent Counter (2)


- Counter with lock is thread safe

```
1  typedef struct __counter_t {
2      int          value;
3      pthread_mutex_t lock;
4  } counter_t;
5
6  void init(counter_t *c) {
7      c->value = 0;
8      Pthread_mutex_init(&c->lock, NULL);
9  }
10
11 void increment(counter_t *c) {
12     Pthread_mutex_lock(&c->lock);
13     c->value++;
14     Pthread_mutex_unlock(&c->lock);
15 }
16
17 void decrement(counter_t *c) {
18     Pthread_mutex_lock(&c->lock);
19     c->value--;
20     Pthread_mutex_unlock(&c->lock);
21 }
22
23 int get(counter_t *c) {
24     Pthread_mutex_lock(&c->lock);
25     int rc = c->value;
26     Pthread_mutex_unlock(&c->lock);
27     return rc;
28 }
```

Figure 29.2: A Counter With Locks

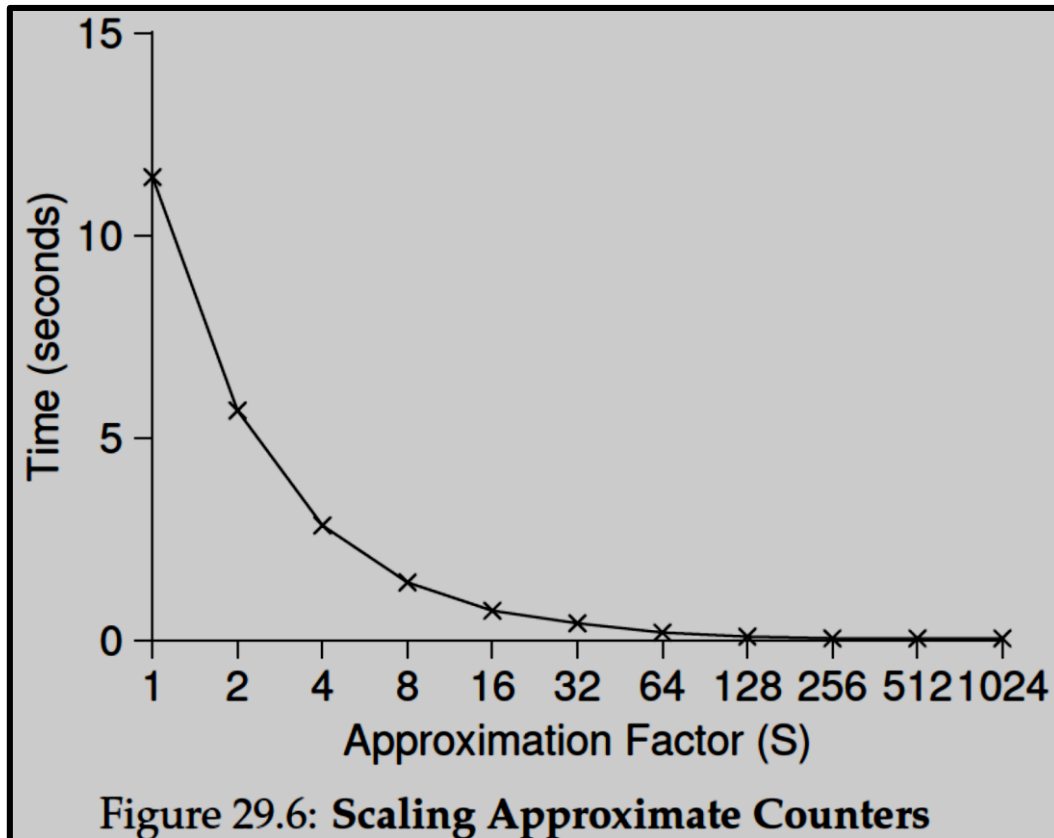
# Concurrent Counter (3)

- Performance of concurrent counter scales poorly
  - Locks add additional overhead which reduces performance
- Solution
  - Use approximate counters which update counter value after some threshold  $S$  so the lock is accessed  $S$  times



# Approximate Counter

- Performance of approximate counter with varying values of  $S$



# Concurrent Linked List

- Per-list lock which is acquired when needed

```
1 // basic node structure
2 typedef struct __node_t {
3     int key;
4     struct __node_t *next;
5 } node_t;
6
7 // basic list structure (one used per list)
8 typedef struct __list_t {
9     node_t *head;
10    pthread_mutex_t lock;
11 } list_t;
12
13 void List_Init(list_t *L) {
14     L->head = NULL;
15     pthread_mutex_init(&L->lock, NULL);
16 }
17
```



# Concurrent Linked List (2)

- Insert Function: Malloc call could be moved out of lock to avoid branching

```
int List_Insert(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        pthread_mutex_unlock(&L->lock);
        return -1; // fail
    }
    new->key = key;
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; // success
}
```

```
void List_Insert(list_t *L, int key) {
    // synchronization not needed
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        return;
    }
    new->key = key;

    // just lock critical section
    pthread_mutex_lock(&L->lock);
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
}
```

# Concurrent Queues

- Implemented using two locks: head and tail locks (to provide concurrency in enqueue and dequeuer operations)

```
typedef struct __node_t {
    int          value;
    struct __node_t *next;
} node_t;

typedef struct __queue_t {
    node_t      *head;
    node_t      *tail;
    pthread_mutex_t head_lock, tail_lock;
}
```

# Concurrent Queues (2)

- Enqueue and Dequeue operations

```
void Queue_Enqueue(queue_t *q, int value) {
    node_t *tmp = malloc(sizeof(node_t));
    assert(tmp != NULL);
    tmp->value = value;
    tmp->next = NULL;

    pthread_mutex_lock(&q->tail_lock);
    q->tail->next = tmp;
    q->tail = tmp;
    pthread_mutex_unlock(&q->tail_lock);
}
```

```
int Queue_Dequeue(queue_t *q, int *value) {
    pthread_mutex_lock(&q->head_lock);
    node_t *tmp = q->head;
    node_t *new_head = tmp->next;
    if (new_head == NULL) {
        pthread_mutex_unlock(&q->head_lock);
        return -1; // queue was empty
    }
    *value = new_head->value;
    q->head = new_head;
    pthread_mutex_unlock(&q->head_lock);
    free(tmp);
    return 0;
}
```

# Concurrent Hash Table

- Made using concurrent linked list
- Uses one lock per hash bucket each represented by a list

```
1 #define BUCKETS (101)
2
3 typedef struct __hash_t {
4     list_t lists[BUCKETS];
5 } hash_t;
6
7 void Hash_Init(hash_t *H) {
8     int i;
9     for (i = 0; i < BUCKETS; i++)
10         List_Init(&H->lists[i]);
11 }
12
13 int Hash_Insert(hash_t *H, int key) {
14     return List_Insert(&H->lists[key % BUCKETS], key);
15 }
16
17 int Hash_Lookup(hash_t *H, int key) {
18     return List_Lookup(&H->lists[key % BUCKETS], key);
19 }
```

Figure 29.10: A Concurrent Hash Table




Figure 29.11: Scaling Hash Tables

# Summary

- We introduced a sampling of concurrent data structures, from counters, to lists and queues, and finally to the ubiquitous and heavily used hash table
- We must be careful with acquisition/release of locks around conditional control flow
- Enabling more concurrency does not necessarily increase performance