

Operating System (OS)

CS232

Concurrency: Condition Variables

Dr. Muhammad Mobeen Movania

Dr. Muhammad Saeed

Condition Variables

- Issues with Mutex Locks
- Condition Variables
- PThread Condition Variables API
- Producer/Consumer problem
- Need of synchronization
- Summary

Example - Lock Issues

```
1  volatile int done = 0;
2
3  void *child(void *arg) {
4      printf("child\n");
5      done = 1;
6      return NULL;
7  }
8
9  int main(int argc, char *argv[]) {
10     printf("parent: begin\n");
11     pthread_t c;
12     Pthread_create(&c, NULL, child, NULL);
13     while (done == 0)
14         ; // spin
15     printf("parent: end\n");
16     return 0;
17 }
```

Issues with Mutex Locks

- Locks provide basic synchronization (mutual exclusion) which does not work in all cases (as discussed in last slide)
 - Spin waiting wastes a lot of CPU cycles
 - What if a thread wanted to wait for a condition?
- Solution
 - Condition variables

Condition Variables

- Are provided by OS Kernel
- Provide two operations
 - wait() – puts a thread in a queue of waiting threads
 - signal() – informs one or more waiting threads
- If a thread wants to be informed of a change in state of a CV
 - it calls wait() on CV and it will be put in a queue associated with that CV
- If a thread wants to inform other threads of a change it CV
 - it calls signal() on CV
 - The OS will awake one (or more) threads waiting in the queue associated with that CV

PThread Condition Variable API

```
pthread_cond_t c;    // initialization
```

```
//wait function
```

```
pthread_cond_wait(pthread_cond_t *c,  
                  pthread_mutex_t *m);
```

```
//signal function
```

```
pthread_cond_signal(pthread_cond_t *c);
```

Example Code

```
pthread_cond_t cv;  
pthread_mutex_t mutex;  
long long balance = 50000000;
```

```
int main (){  
    pthread_t tid;  
    pthread_cond_init(&cv, NULL);  
    pthread_mutex_init(&mutex, NULL);  
    pthread_create (&tid, NULL, debit, NULL);  
  
    pthread_mutex_lock(&mutex);  
    if (balance != 0)  
        pthread_cond_wait(&cv, &mutex);  
    pthread_mutex_unlock(&mutex);  
  
    printf("balance = %lld \n", balance);  
    pthread_mutex_destroy(&mutex);  
    pthread_cond_destroy(&cv);  
    return 0;  
}
```

```
void* debit(void* arg){  
    while(balance > 0){  
        pthread_mutex_lock(&mutex);  
        balance--;  
        pthread_mutex_unlock(&mutex);  
    }  
    pthread_cond_signal(&cv);  
    return NULL;  
}
```

```
balance = 0
```

Producer-Consumer

- A classic synchronization problem
- One (or more) Producer thread(s) put values in a buffer
- One (or more) Consumer threads(s) consume values from the buffer
- The buffer is bounded
- How do we synchronize them?

Needs synchronization

```
1  int buffer;
2  int count = 0; // initially, empty
3
4  void put(int value) {
5      assert(count == 0);
6      count = 1;
7      buffer = value;
8  }
9
10 int get() {
11     assert(count == 1);
12     count = 0;
13     return buffer;
14 }
```

```
1  void *producer(void *arg) {
2      int i;
3      int loops = (int) arg;
4      for (i = 0; i < loops; i++) {
5          put(i);
6      }
7  }
8
9  void *consumer(void *arg) {
10     int i;
11     while (1) {
12         int tmp = get();
13         printf("%d\n", tmp);
14     }
15 }
```

Needs synchronization ... contd.

- Problems?

```
1  int loops; // must initialize somewhere...
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);           // p1
9          if (count == 1)                       // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11             put(i);                             // p4
12             Pthread_cond_signal(&cond);         // p5
13             Pthread_mutex_unlock(&mutex);       // p6
14         }
15     }
16
17     void *consumer(void *arg) {
18         int i;
19         for (i = 0; i < loops; i++) {
20             Pthread_mutex_lock(&mutex);         // c1
21             if (count == 0)                     // c2
22                 Pthread_cond_wait(&cond, &mutex); // c3
23             int tmp = get();                     // c4
24             Pthread_cond_signal(&cond);         // c5
25             Pthread_mutex_unlock(&mutex);       // c6
26             printf("%d\n", tmp);
27         }
28     }
```

Example Run (c1,p,c2,c1)

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
c1	Running		Ready		Ready	0	Nothing to get
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	Buffer now full T_{c1} awoken
	Sleep		Ready	p4	Running	1	
	Ready		Ready	p5	Running	1	
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	Buffer full; sleep T_{c2} sneaks in ...
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	
	Ready	c1	Running		Sleep	1	
	Ready	c2	Running		Sleep	1	... and grabs data T_p awoken
	Ready	c4	Running		Sleep	0	
	Ready	c5	Running		Ready	0	
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

Needs synchronization ... contd.

- Scenario
 - 1 Producer
 - 2 Consumers
- C0 runs, sleeps.
- P0 runs, produces, sleeps.
- C1 runs, consumes signals the CV.
- Which thread is woken?

```
1  int loops;
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);          // p1
9          while (count == 1)                    // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11             put(i);                            // p4
12             Pthread_cond_signal(&cond);        // p5
13             Pthread_mutex_unlock(&mutex);      // p6
14         }
15     }
16
17 void *consumer(void *arg) {
18     int i;
19     for (i = 0; i < loops; i++) {
20         Pthread_mutex_lock(&mutex);          // c1
21         while (count == 0)                    // c2
22             Pthread_cond_wait(&cond, &mutex); // c3
23         int tmp = get();                      // c4
24         Pthread_cond_signal(&cond);          // c5
25         Pthread_mutex_unlock(&mutex);        // c6
26         printf("%d\n", tmp);
27     }
28 }
```

Example Run (c1,c2,p,c1,c2)

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
c1	Running		Ready		Ready	0	Nothing to get
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	
	Sleep	c1	Running		Ready	0	Nothing to get
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	
	Sleep		Sleep	p1	Running	0	Buffer now full T_{c1} awoken
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	
	Ready		Sleep	p5	Running	1	Must sleep (full) Recheck condition T_{c1} grabs data Oops! Woke T_{c2}
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	Nothing to get
	Ready		Sleep	p3	Sleep	1	
c2	Running		Sleep		Sleep	1	
c4	Running		Sleep		Sleep	0	Nothing to get
c5	Running		Ready		Sleep	0	
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	Everyone asleep...
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	
	Sleep	c2	Running		Sleep	0	Everyone asleep...
	Sleep	c3	Sleep		Sleep	0	

Use 2 CVs

- Solution
 - Use 2 CVs

```
1  cond_t  empty, fill;
2  mutex_t mutex;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          Pthread_mutex_lock(&mutex);
8          while (count == 1)
9              Pthread_cond_wait(&empty, &mutex);
10         put(i);
11         Pthread_cond_signal(&fill);
12         Pthread_mutex_unlock(&mutex);
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         Pthread_mutex_lock(&mutex);
20         while (count == 0)
21             Pthread_cond_wait(&fill, &mutex);
22         int tmp = get();
23         Pthread_cond_signal(&empty);
24         Pthread_mutex_unlock(&mutex);
25         printf("%d\n", tmp);
26     }
27 }
```

Multiple buffers

```
1  int buffer[MAX];
2  int fill_ptr = 0;
3  int use_ptr  = 0;
4  int count    = 0;
5
6  void put(int value) {
7      buffer[fill_ptr] = value;
8      fill_ptr = (fill_ptr + 1) % MAX;
9      count++;
10 }
11
12 int get() {
13     int tmp = buffer[use_ptr];
14     use_ptr = (use_ptr + 1) % MAX;
15     count--;
16     return tmp;
17 }
```

```
1  cond_t empty, fill;
2  mutex_t mutex;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          Pthread_mutex_lock(&mutex);
8          while (count == MAX)
9              Pthread_cond_wait(&empty, &mutex);
10         put(i);
11         Pthread_cond_signal(&fill);
12         Pthread_mutex_unlock(&mutex);
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         Pthread_mutex_lock(&mutex);
20         while (count == 0)
21             Pthread_cond_wait(&fill, &mutex);
22         int tmp = get();
23         Pthread_cond_signal(&empty);
24         Pthread_mutex_unlock(&mutex);
25         printf("%d\n", tmp);
26     }
27 }
```


Covering conditions

- Solution
 - Broadcast signal

```
1 // how many bytes of the heap are free?
2 int bytesLeft = MAX_HEAP_SIZE;
3
4 // need lock and condition too
5 cond_t c;
6 mutex_t m;
7
8 void *
9 allocate(int size) {
10     Pthread_mutex_lock(&m);
11     while (bytesLeft < size)
12         Pthread_cond_wait(&c, &m);
13     void *ptr = ...; // get mem from heap
14     bytesLeft -= size;
15     Pthread_mutex_unlock(&m);
16     return ptr;
17 }
18
19 void free(void *ptr, int size) {
20     Pthread_mutex_lock(&m);
21     bytesLeft += size;
22     Pthread_cond_signal(&c); // whom to signal??
23     Pthread_mutex_unlock(&m);
24 }
```


Summary

- We have seen a new synchronization primitive called a condition variable
- Unlike locks, condition variable does not spin wait, instead, it provides two function wait and signal
 - wait() call puts the calling thread in a wait queue associated with the condition variable
 - signal() call awakes one or more waiting thread in the condition variable waiting queue
- Condition variables enable us to neatly solve producer/consumer problem, as well as covering conditions.