ⓘ Students have either already taken or started taking this quiz, so take care when editing it. If you change any quiz questions in a significant way, you might want to consider re-grading students' quizzes who took the old version of the quiz.

Points 8 Published ⋮

Details Questions

Show question details

Question 1 pts

The following code is the implementation of a lock using load-linked/store-conditional instructions:

```
1  void lock(lock_t *lock) {
2      while (1) {
3          while (LoadLinked(&lock->flag) == 1)
4              ; // spin until it's zero
5          if (StoreConditional(&lock->flag, 1) == 1)
6              return; // if set-it-to-1 was a success: all done
7                      // otherwise: try it all over again
8      }
9  }
10
11 void unlock(lock_t *lock) {
12     lock->flag = 0;
13 }
```

Which of the following is most correct about this lock?

- "If" condition (on line 5) will be TRUE if no other thread called store-conditional on this lock since this thread's "while" loop (on line 3) ended
- "If" condition (on line 5) will be FALSE if no other thread called store-conditional on this lock since this thread's "while" loop (on line 3) ended
- "If" condition (on line 5) will be TRUE if no other thread called load-linked on this lock since this thread's "while" loop (on line 3) ended
- "If" condition (on line 5) will be FALSE if no other thread called load-linked on this lock since this thread's "while" loop (on line 3) ended

Question 1 pts

The following is the implementation of a lock using fetch-and-add instructions:

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void lock_init(lock_t *lock) {
7      lock->ticket = 0;
8      lock->turn   = 0;
9  }
10
11 void lock(lock_t *lock) {
12     int myturn = FetchAndAdd(&lock->ticket);
13     while (lock->turn != myturn)
14         ; // spin
15  }
16
17 void unlock(lock_t *lock) {
18     lock->turn = lock->turn + 1;
19 }
```

```
1  int FetchAndAdd(int *ptr) {
2      int old = *ptr;
3      *ptr = old + 1;
4      return old;
5  }
```

Figure 28.7: **Ticket Locks**

Suppose the current value of ticket = 2 and the current value of turn = 1.

A thread calls lock(). Then:

- This thread will spin until another thread calls unlock()
- This thread will spin until two other threads call unlock()
- This thread will spin until three other threads call unlock()
- This thread will not spin at all

Question 1 pts

Here is the implementation of a lock using yeild() function:

```
1  void init() {
2      flag = 0;
3  }
4
5  void lock() {
6      while (TestAndSet(&flag, 1) == 1)
7          yield(); // give up the CPU
8  }
9
10 void unlock() {
11     flag = 0;
12 }
```

Suppose the lock was acquired by another thread when lock() is called by a thread. Then

- This thread keeps getting scheduled by the OS, but the thread keeps giving up the CPU
- This thread goes to sleep is woken up only when the lock is released
- This thread keeps spinning until the lock is released by the other thread

This thread gets scheduled by the OS once, and it gives up the CPU. The thread is scheduled again only after lock is released by the other thread.

Question 1 pts

Here is the code for a queue-based implementation of a lock:

```
typedef struct __lock_t {
    int flag;
    int guard;
    queue_t *q;
} lock_t;

void lock_init(lock_t *m) {
    m->flag  = 0;
    m->guard = 0;
    queue_init(m->q);
}
```

```
void lock(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (m->flag == 0) {
        m->flag = 1; // lock is acquired
        m->guard = 0;
    } else {
        queue_add(m->q, gettid());
        m->guard = 0;
        park();
    }
}

void unlock(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (queue_empty(m->q))
        m->flag = 0; // let go of lock; no one wants it
    else
        unpark(queue_remove(m->q)); // hold lock
                                    // (for next thread!)
    m->guard = 0;
}
```

Which of the following is most accurate about this implementation?

- "guard" is a spin lock around actual the actual locking mechanisms of "flag" and "queue"
- "flag" is a spin lock around actual the actual locking mechanisms of "guard" and "queue"
- Both "flag" and "guard" serve as the main locking mechanism and there is no spinning
- Both "flag" and "guard" are spin-locks

Question 1 pts

Here is the code for queue-based lock implementation:

```
13  void lock(lock_t *m) {
14      while (TestAndSet(&m->guard, 1) == 1)
15          ; //acquire guard lock by spinning
16      if (m->flag == 0) {
17          m->flag = 1; // lock is acquired
18          m->guard = 0;
19      } else {
20          queue_add(m->q, gettid());
21          m->guard = 0;
22          park();
23      }
24  }
25
26  void unlock(lock_t *m) {
27      while (TestAndSet(&m->guard, 1) == 1)
28          ; //acquire guard lock by spinning
29      if (queue_empty(m->q))
30          m->flag = 0; // let go of lock; no one wants it
31      else
32          unpark(queue_remove(m->q)); // hold lock
33                                      // (for next thread!)
34      m->guard = 0;
35  }
```

```
1   typedef struct __lock_t {
2       int flag;
3       int guard;
4       queue_t *q;
5   } lock_t;
6
7   void lock_init(lock_t *m) {
8       m->flag  = 0;
9       m->guard = 0;
10      queue_init(m->q);
11  }
```

Suppose a thread that had acquired the lock calls unlock(), while another thread was waiting in the queue.

Then which of the following is most accurate:

- "flag" was set to 1, and its value is unchanged as the waiting thread wakes up and acquires the lock.
- "flag" was set to 0, and its value is unchanged as the waiting thread wakes up and acquires the lock.
- "flag" was set to 1, and its value changes to 0 as the lock is released. It is then set to 1 after the waiting thread wakes up and acquires the lock.
- "flag" was set to 0, and its value changes to 1 as the lock is released. It is then set to 0 after the waiting thread wakes up and acquires the lock.

Question 1 pts

Which of the following statements best describes concurrent approximate counters?

- Concurrent approximate counters trade off exact counting accuracy for improved performance in parallel and concurrent environments.
- Concurrent approximate counters guarantee exact counting accuracy in multi-threaded environments.
- Concurrent approximate counters are designed to provide precise counting results even when multiple threads access the counter simultaneously.
- Concurrent approximate counters are exclusively used in single-threaded applications to avoid counting errors.

Question 1 pts

Which of the following best describes the purpose of the following C code snippet?

```
if (raise(SIGUSR1) != 0)
    perror("Failed to raise SIGUSR1");
```

- It causes a process to send the SIGUSR1 signal to itself
- The code increases the value of SIGUSR1 in the current process.
- The code is used to catch and handle the SIGUSR1 signal in the current process.

It creates a new process and sends the SIGUSR1 signal to it.

Question 1 pts

What is the purpose of the following C code?

```
#include <unistd.h>

int main(void) {
    alarm(10);
    for (;;) {
        // Infinite loop
    }
}
```

ıswer

- The program sets an alarm to trigger after 10 seconds and enters an infinite loop.
- The code sets an alarm for 10 seconds and then terminates the program.
- It creates a background process that runs for 10 seconds and then terminates.
- It waits for 10 seconds and then prints a message to the console.

+ New question + New question group Find questions

Notify users this quiz has changed

Cancel Save