

# Deep\_Learning\_Lab\_Exam\_Unsolved\_Fall\_2021

November 8, 2024

## 1 Deep Learning Lab Exam

In this lab exam, you will be implementing a neural network with 1 hidden layer (with the activation function tanh) and use it to classify whether a banknote is authentic (1) or not (0): Your network would look like the following:

## 2 Exam set up (Do not edit!)

```
[1]: # Download the dataset using wget. If you are not on colab or linux, you can
      ↪ skip this line and download your dataset directly
      # and place it in the same folder as your notebook
      !wget http://archive.ics.uci.edu/ml/machine-learning-databases/00267/
      ↪ data_banknote_authentication.txt
```

```
--2024-11-08 09:55:15-- http://archive.ics.uci.edu/ml/machine-learning-
databases/00267/data_banknote_authentication.txt
Resolving archive.ics.uci.edu (archive.ics.uci.edu)...
128.195.10.252
Connecting to archive.ics.uci.edu (archive.ics.uci.edu)|128.195.10.252|:80...
connected.
HTTP request sent, awaiting response... 200 OK
Length: unspecified
Saving to: 'data_banknote_authentication.txt'

data_banknote_auth     [  <=>          ] 45.31K  76.2KB/s   in 0.6s

2024-11-08 09:55:16 (76.2 KB/s) - 'data_banknote_authentication.txt' saved
[46400]
```

```
[25]: # Import relevant libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import pandas as pd
```

Dataset Attribute Information:

1. variance of Wavelet Transformed image (continuous)
2. skewness of Wavelet Transformed image (continuous)
3. curtosis of Wavelet Transformed image (continuous)
4. entropy of image (continuous)
5. class (integer)

```
[26]: # Importing the dataset and extracting the inputs and the outputs
data = np.genfromtxt("data_banknote_authentication.txt", delimiter = ",")
X = data[:, :4]
y = data[:, 4]
```

```
[27]: # Split the dataset into a training and test dataset

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
X_train = X_train.T
y_train = y_train.reshape(1, y_train.shape[0])
X_test = X_test.T
y_test = y_test.reshape(1, y_test.shape[0])
print ("Train X Shape: ", X_train.shape)
print ("Train Y Shape: ", y_train.shape)
print ("%d training examples." % (X_train.shape[1]))

print ('\nTest X Shape: ', X_test.shape)
```

Train X Shape: (4, 1097)

Train Y Shape: (1, 1097)

1097 training examples.

Test X Shape: (4, 275)

Note that in X, every column contains each record

### 3 The exam starts

In this exam, we will: 1. Define the neural network structure ( # of input neurons, # of hidden neurons, etc). 2. Initialize the model's parameters 3. Loop: - Implement forward propagation - Compute loss - Implement backward propagation to get the gradients - Update parameters (gradient descent) 4. Evaluate accuracy

#### 3.0.1 [1 mark] Define Structure:

In this task, you need to take the inputs (**X**) and output(**Y**) and number of hidden neurons you want to add in your hidden layer, and based on that return number of input neurons, number of hidden neurons and number of output neurons.

```
[28]: def define_structure(X, Y,hidden_neurons):
    # IMPLEMENT
    num_input_neurons = 4
    num_hidden_neurons = hidden_neurons
    num_output_neurons = 1
    return num_input_neurons, num_hidden_neurons, num_output_neurons

num_input_neurons, num_hidden_neurons, num_output_neurons = define_structure(X_train, y_train,4)
print("The size of the input layer is := " + str(num_input_neurons))
print("The size of the hidden layer is := " + str(num_hidden_neurons))
print("The size of the output layer is := " + str(num_output_neurons))

# Note that input layer should always have 4 neurons since we have 1 input
# Output layer should have 1 neuron since we are predicting 1 value, whether
# the note is authentic or not

# For the purpose of this exam, number of hidden neurons would be 4.
```

The size of the input layer is := 4  
The size of the hidden layer is := 4  
The size of the output layer is := 1

### 3.0.2 [1 mark] : Parameter Initialization

Now that we know the structure of the network, we can initialize the parameters of our network.

Since our network has 1 hidden layer, we will have 2 set of weights and biases. 1. From input to hidden layer (W1, b1) 2. from Hidden to output layer (W2, b2)

Based on the structure of the network, initialize these parameters.

Initialize weights (W1,W2) using `np.random.randn` and scale them by a factor of 0.01.

Initialize biases (b1,b2) as 0s (using `np.zeros`)

```
[29]: def parameters_initialization(num_input_neurons, num_hidden_neurons, num_output_neurons):
    # Seed is set for ease of testing. Do not remove.
    np.random.seed(2)
    # IMPLEMENT

    W1 = np.random.randn(num_hidden_neurons, num_input_neurons) * 0.01
    b1 = np.zeros((num_hidden_neurons, 1))

    W2 = np.random.randn(num_output_neurons, num_hidden_neurons) * 0.01
    b2 = np.zeros((num_output_neurons, 1))

    # We are going to be storing the parameters in a
```

```
# dictionary for ease of access later on
parameters = {"W1": W1,
              "b1": b1,
              "W2": W2,
              "b2": b2}

return parameters
```

### 3.0.3 [0.5 marks]: Implement the Sigmoid function

Sigmoid is defined as follows:

$$S(z) = \frac{1}{1 + e^{-z}}$$

```
[30]: def sigmoid(z):
      # IMPLEMENT
      denom = 1 + np.exp(-z)
      sig = np.divide(1, denom)
      return sig
```

### 3.0.4 [1 mark] Forward propogation

In order to get a result from the neural network, we need to propogate our inputs through the network.

Your goal is to implement a function which propogates the inputs (X) through the network and computes the result at each part of the network.

We are naming the parts of the networks as follows:

$Z_1$  as the output to the hidden layer from the input layer (without the activation function applied)

$A_1$  as applying the activation function to  $Z_1$

$Z_2$  as the output to the output layer from the hidden layer ( $A_1$ ) with no activation function

$A_2$  as sigmoid being applied to  $Z_2$  to get the final output (Probability that the input belongs to class 1 by the network)

$$Z_1 = W_1 X + b_1$$

$$A_1 = \tanh(Z_1)$$

$$Z_2 = W_2 A_1 + b_2$$

$$A_2 = S(Z_2)$$

**Note:** Make sure your matrices are in such a way that your weight matrices are initialized in such a way, such that these equations make sense. If you don't, then you might have to adjust how you deal with backpropogation.

In numpy, you can do tanh, by `np.tanh`

```
[31]: def forward_propagation(X, parameters):
    # Extract weights and biases from the dictionary
    W1 = parameters['W1']
    b1 = parameters['b1']
    W2 = parameters['W2']
    b2 = parameters['b2']

    # Compute Z1,A1,Z2,A2
    # IMPLEMENT

    Z1 = np.dot(W1, X) + b1
    A1 = np.tanh(Z1)
    Z2 = np.dot(W2, A1) + b2
    A2 = sigmoid(Z2)

    # Save the output of each layer in a dictionary to help with backpropagation,
    ↪ later on
    cache = {"Z1": Z1, "A1": A1, "Z2": Z2, "A2": A2}
    # Return the final output and output at each layer
    return A2, cache
```

### 3.0.5 [1.5 marks]: Binary Cross-entropy loss

In order for you to evaluate how far your model is from correct predictions, you need to have a loss function. We will be using Binary Cross-entropy loss since our prediction can be either Class 0 or Class 1. The formula to compute Binary Cross-entropy loss  $J$  using actual outputs  $Y$  and the output from your last layer  $A^{[2]}$  is as follows

$$J(Y, A^{[2]}) = -\frac{1}{m} \sum_{i=1}^m Y^{(i)} \log(A^{[2](i)}) + (1 - Y^{(i)}) \log(1 - A^{[2](i)})$$

where  $m$  is the number of samples.

You can use `np.log` for log.

```
[36]: def cross_entropy_cost(Y,A2):
    # IMPLEMENT
    p1 = np.multiply(Y, np.log(A2)); p2 = np.multiply(1 - Y, np.log(1 - A2))
    summ = np.sum(p1 + p2)
    cost = np.multiply(np.divide(-1, Y.shape[1]), summ)
    return cost
```

### 3.0.6 [1.5 marks] : Back Propagation

Now that we have gotten the output at each layer, we can compute the gradients at each layer and use that in order to nudge our model parameters in the right direction.

To figure out how much our actual output deviates from our desired output.

$$dZ^{[2]} = A^{[2]} - Y$$

and using that we can compute the changes for the weights and biases in the hidden layer

$$dW^{[2]} = \frac{1}{m} dZ^{[2]} A^{[1]T}$$

$db^{[2]}$  is what you can get if you sum each row for  $dZ^{[2]}$  and then divide this sum by number of samples  $m$  (Do keep dimensions in numpy!)

To compute the changes in the input layer weights

$$dZ^{[1]} = W^{[2]T} dZ^{[2]} * g^{[1]'}(Z^{[1]})$$

Where  $g^{[1]}$  is your activation function (in our case is tanh)

The derivative of tanh is

$$\tanh(Z) = 1 - \tanh^2(Z)$$

Note that  $*$  means element-wise multiplication. In numpy, for matrix multiplication use `np.dot`, and for elementwise you can use `*`.

Using that, we can compute the change needed in weights and biases.

$$dW^{[1]} = \frac{1}{m} dZ^{[1]} X^T$$

$db^{[1]}$  is what you can get if you sum each row for  $dZ^{[1]}$  and then divide this sum by number of samples  $m$  (Do keep dimensions in numpy!)

Note: sum each row for means that if I have an array `[[1,2],[3,5]]`, we will get `[3,8]` and keeping the dimensions same, would give us `[[3],[8]]`.

To raise each element in a numpy array A to the power i, you can can be done by `np.power(A,i)`

```
[37]: def backward_propagation(parameters, cache, X, Y):

    W1 = parameters['W1']
    W2 = parameters['W2']
    A1 = cache['A1']
    A2 = cache['A2']

    # Compute dW2, db2, dW1, db1
    # You would find it convenient it to compute dZ2, dZ1
    # IMPLEMENT
    Z1 = cache['Z1']
    dZ2 = A2 - Y
    dW2 = np.multiply(np.divide(1, Y.shape[1]), np.dot(dZ2, A1.T))
    dZ1 = np.multiply(np.dot(W2.T, dZ2), (1 - np.tanh(Z1)**2))
    dW1 = np.multiply(np.divide(1, Y.shape[1]), np.dot(dZ1, X.T))
```

```

    db2 = np.multiply(np.divide(1, Y.shape[1]), np.sum(dZ2, axis=1,
↪keepdims=True))
    db1 = np.multiply(np.divide(1, Y.shape[1]), np.sum(dZ1, axis=1,
↪keepdims=True))

    # Store the changes in a dictionary grads to be used for gradient descent
    grads = {"dW1": dW1, "db1": db1, "dW2": dW2, "db2": db2}

    return grads

```

### 3.0.7 [1.5 marks]: Gradient Descent

Now that we have computed the changes that we need, we can just apply gradient descent to every weight and bias parameter with learning rate  $\alpha$

$$W = W - \alpha dW \text{ and } B = b - \alpha db$$

```

[38]: def gradient_descent(parameters, grads, learning_rate = 0.01):
    W1 = parameters['W1']
    b1 = parameters['b1']
    W2 = parameters['W2']
    b2 = parameters['b2']

    dW1 = grads['dW1']
    db1 = grads['db1']
    dW2 = grads['dW2']
    db2 = grads['db2']

    # Perform gradient descent
    # IMPLEMENT

    W1 = W1 - np.multiply(learning_rate, dW1)
    b1 = b1 - np.multiply(learning_rate, db1)
    W2 = W2 - np.multiply(learning_rate, dW2)
    b2 = b2 - np.multiply(learning_rate, db2)

    # Return the updated parameters
    parameters = {"W1": W1, "b1": b1, "W2": W2, "b2": b2}

    return parameters

```

### 3.0.8 [1.5 marks] Making the network: Bringing it all together.

Now we can bring everything together and make our model work.

As a reminder, the network is supposed to

1. Define the neural network structure ( # of input neurons, # of hidden neurons, etc).

2. Initialize the model's parameters
3. Loop:
  - Implement forward propagation
  - Compute loss (print the loss)
  - Implement backward propagation to get the gradients
  - Update parameters (gradient descent)

Once the model is trained, we can return the parameters.

```
[39]: def neural_network_model(X, Y, num_hidden_neurons, num_iterations = 1000):
    np.random.seed(3)
    # Define the structure
    # Initialize Parameters
    # IMPLEMENT

    num_input_neurons, num_hidden_neurons, num_output_neurons =
    define_structure(X, Y, num_hidden_neurons)
    parameters = parameters_initialization(num_input_neurons,
    num_hidden_neurons, num_output_neurons)

    # Training loop

    for i in range(num_iterations):
        # forward prop
        A2, cache = forward_propagation(X, parameters)
        # Compute loss
        loss = cross_entropy_cost(Y, A2)
        # Back prop to find gradients
        grads = backward_propagation(parameters, cache, X, Y)
        # Update gradients
        parameters = gradient_descent(parameters, grads)
        # IMPLEMENT
        # Print loss (Do not edit!)
        if i % 5 == 0:
            print("Cost after iteration %i: %f" % (i, loss))
    return parameters

# Train the model (Do not edit!!)
parameters = neural_network_model(X_train, y_train, 4, num_iterations=1000)
```

```
Cost after iteration 0: 0.692975
Cost after iteration 5: 0.692665
Cost after iteration 10: 0.692351
Cost after iteration 15: 0.692028
Cost after iteration 20: 0.691692
Cost after iteration 25: 0.691338
Cost after iteration 30: 0.690962
Cost after iteration 35: 0.690557
```



Cost after iteration 40: 0.690117  
Cost after iteration 45: 0.689636  
Cost after iteration 50: 0.689107  
Cost after iteration 55: 0.688521  
Cost after iteration 60: 0.687873  
Cost after iteration 65: 0.687152  
Cost after iteration 70: 0.686352  
Cost after iteration 75: 0.685463  
Cost after iteration 80: 0.684477  
Cost after iteration 85: 0.683388  
Cost after iteration 90: 0.682187  
Cost after iteration 95: 0.680866  
Cost after iteration 100: 0.679420  
Cost after iteration 105: 0.677843  
Cost after iteration 110: 0.676127  
Cost after iteration 115: 0.674268  
Cost after iteration 120: 0.672259  
Cost after iteration 125: 0.670096  
Cost after iteration 130: 0.667773  
Cost after iteration 135: 0.665286  
Cost after iteration 140: 0.662629  
Cost after iteration 145: 0.659797  
Cost after iteration 150: 0.656787  
Cost after iteration 155: 0.653594  
Cost after iteration 160: 0.650215  
Cost after iteration 165: 0.646648  
Cost after iteration 170: 0.642889  
Cost after iteration 175: 0.638938  
Cost after iteration 180: 0.634794  
Cost after iteration 185: 0.630457  
Cost after iteration 190: 0.625929  
Cost after iteration 195: 0.621210  
Cost after iteration 200: 0.616305  
Cost after iteration 205: 0.611217  
Cost after iteration 210: 0.605950  
Cost after iteration 215: 0.600512  
Cost after iteration 220: 0.594907  
Cost after iteration 225: 0.589145  
Cost after iteration 230: 0.583233  
Cost after iteration 235: 0.577180  
Cost after iteration 240: 0.570995  
Cost after iteration 245: 0.564687  
Cost after iteration 250: 0.558268  
Cost after iteration 255: 0.551746  
Cost after iteration 260: 0.545132  
Cost after iteration 265: 0.538436  
Cost after iteration 270: 0.531668  
Cost after iteration 275: 0.524840

Cost after iteration 280: 0.517962  
Cost after iteration 285: 0.511044  
Cost after iteration 290: 0.504098  
Cost after iteration 295: 0.497134  
Cost after iteration 300: 0.490163  
Cost after iteration 305: 0.483197  
Cost after iteration 310: 0.476244  
Cost after iteration 315: 0.469317  
Cost after iteration 320: 0.462425  
Cost after iteration 325: 0.455576  
Cost after iteration 330: 0.448781  
Cost after iteration 335: 0.442047  
Cost after iteration 340: 0.435381  
Cost after iteration 345: 0.428791  
Cost after iteration 350: 0.422283  
Cost after iteration 355: 0.415862  
Cost after iteration 360: 0.409532  
Cost after iteration 365: 0.403299  
Cost after iteration 370: 0.397165  
Cost after iteration 375: 0.391133  
Cost after iteration 380: 0.385205  
Cost after iteration 385: 0.379383  
Cost after iteration 390: 0.373667  
Cost after iteration 395: 0.368060  
Cost after iteration 400: 0.362561  
Cost after iteration 405: 0.357170  
Cost after iteration 410: 0.351887  
Cost after iteration 415: 0.346712  
Cost after iteration 420: 0.341643  
Cost after iteration 425: 0.336679  
Cost after iteration 430: 0.331820  
Cost after iteration 435: 0.327064  
Cost after iteration 440: 0.322409  
Cost after iteration 445: 0.317854  
Cost after iteration 450: 0.313397  
Cost after iteration 455: 0.309037  
Cost after iteration 460: 0.304771  
Cost after iteration 465: 0.300598  
Cost after iteration 470: 0.296516  
Cost after iteration 475: 0.292523  
Cost after iteration 480: 0.288616  
Cost after iteration 485: 0.284795  
Cost after iteration 490: 0.281056  
Cost after iteration 495: 0.277399  
Cost after iteration 500: 0.273820  
Cost after iteration 505: 0.270319  
Cost after iteration 510: 0.266893  
Cost after iteration 515: 0.263540

Cost after iteration 520: 0.260259  
Cost after iteration 525: 0.257048  
Cost after iteration 530: 0.253905  
Cost after iteration 535: 0.250828  
Cost after iteration 540: 0.247815  
Cost after iteration 545: 0.244866  
Cost after iteration 550: 0.241977  
Cost after iteration 555: 0.239148  
Cost after iteration 560: 0.236378  
Cost after iteration 565: 0.233663  
Cost after iteration 570: 0.231004  
Cost after iteration 575: 0.228398  
Cost after iteration 580: 0.225845  
Cost after iteration 585: 0.223343  
Cost after iteration 590: 0.220889  
Cost after iteration 595: 0.218485  
Cost after iteration 600: 0.216127  
Cost after iteration 605: 0.213815  
Cost after iteration 610: 0.211547  
Cost after iteration 615: 0.209323  
Cost after iteration 620: 0.207141  
Cost after iteration 625: 0.205000  
Cost after iteration 630: 0.202900  
Cost after iteration 635: 0.200838  
Cost after iteration 640: 0.198815  
Cost after iteration 645: 0.196829  
Cost after iteration 650: 0.194879  
Cost after iteration 655: 0.192964  
Cost after iteration 660: 0.191084  
Cost after iteration 665: 0.189237  
Cost after iteration 670: 0.187423  
Cost after iteration 675: 0.185641  
Cost after iteration 680: 0.183890  
Cost after iteration 685: 0.182170  
Cost after iteration 690: 0.180479  
Cost after iteration 695: 0.178817  
Cost after iteration 700: 0.177183  
Cost after iteration 705: 0.175577  
Cost after iteration 710: 0.173998  
Cost after iteration 715: 0.172445  
Cost after iteration 720: 0.170918  
Cost after iteration 725: 0.169415  
Cost after iteration 730: 0.167938  
Cost after iteration 735: 0.166484  
Cost after iteration 740: 0.165054  
Cost after iteration 745: 0.163646  
Cost after iteration 750: 0.162261  
Cost after iteration 755: 0.160898

Cost after iteration 760: 0.159556  
Cost after iteration 765: 0.158234  
Cost after iteration 770: 0.156934  
Cost after iteration 775: 0.155653  
Cost after iteration 780: 0.154392  
Cost after iteration 785: 0.153150  
Cost after iteration 790: 0.151926  
Cost after iteration 795: 0.150721  
Cost after iteration 800: 0.149534  
Cost after iteration 805: 0.148364  
Cost after iteration 810: 0.147212  
Cost after iteration 815: 0.146076  
Cost after iteration 820: 0.144957  
Cost after iteration 825: 0.143854  
Cost after iteration 830: 0.142767  
Cost after iteration 835: 0.141695  
Cost after iteration 840: 0.140639  
Cost after iteration 845: 0.139597  
Cost after iteration 850: 0.138570  
Cost after iteration 855: 0.137557  
Cost after iteration 860: 0.136558  
Cost after iteration 865: 0.135573  
Cost after iteration 870: 0.134602  
Cost after iteration 875: 0.133643  
Cost after iteration 880: 0.132698  
Cost after iteration 885: 0.131765  
Cost after iteration 890: 0.130845  
Cost after iteration 895: 0.129937  
Cost after iteration 900: 0.129041  
Cost after iteration 905: 0.128156  
Cost after iteration 910: 0.127284  
Cost after iteration 915: 0.126422  
Cost after iteration 920: 0.125572  
Cost after iteration 925: 0.124733  
Cost after iteration 930: 0.123904  
Cost after iteration 935: 0.123086  
Cost after iteration 940: 0.122279  
Cost after iteration 945: 0.121481  
Cost after iteration 950: 0.120694  
Cost after iteration 955: 0.119916  
Cost after iteration 960: 0.119148  
Cost after iteration 965: 0.118390  
Cost after iteration 970: 0.117641  
Cost after iteration 975: 0.116901  
Cost after iteration 980: 0.116170  
Cost after iteration 985: 0.115447  
Cost after iteration 990: 0.114734  
Cost after iteration 995: 0.114029

```
[ ]:
```

### 3.0.9 [0.5 marks] Prediction and accuracy:

We now have a well trained model. We can now evaluate how good our model is on the test dataset, your goal is to create a function `prediction` which takes in parameters and the test inputs and based on that give predictions.

If we have probability of class 1 as  $> 0.5$ , our prediction is 1, else 0.

Once you are done with that compute the accuracy: See the **percentage** of predictions that were right (This will be done using actual Ys and your predictions) on both the **training and test** dataset

```
[44]: def prediction(parameters, X):  
    # IMPLEMENT  
    A2, cache = forward_propagation(X, parameters)  
  
    prediction = (A2 > 0.5).astype(int)  
    return prediction  
  
    # Feel free to make a function for accuracy as well
```

```
[45]: predictions = prediction(parameters, X_train)  
    # Feel free to make an accuracy function to compute accuracy  
    predictions = prediction(parameters, X_test)
```

```
[46]: predictions
```

```
[46]: array([[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0,  
          1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0,  
          1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,  
          0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1,  
          1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1,  
          1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 0,  
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0,  
          1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
          1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0,  
          0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1,  
          1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1,  
          0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,  
          1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0]])
```