

# RQ3: DL Refresher

- Due Jan 31 at 4:59pm
- Points 4
- Questions 10
- Time Limit 15 Minutes
- Allowed Attempts Unlimited

Take the Quiz Again

## Attempt History

	Attempt	Time	Score
LATEST	<a href="#">Attempt 1</a>	6 minutes	3.6 out of 4

❗ Correct answers are hidden.

Score for this attempt: 3.6 out of 4

Submitted Jan 31 at 4:43pm

This attempt took 6 minutes.



Question 1

0.4 / 0.4 pts

Consider this simple PyTorch model for regression:

```
import torch
import torch.nn as nn

class SimpleMLP(nn.Module):
    def __init__(self):
        super(SimpleMLP, self).__init__()
        self.fc1 = nn.Linear(4, 3)  # from 4 inputs to 3 neurons
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(3, 1)  # from 3 neurons to 1 output

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        return x
```

How many trainable parameters (weights + biases) does this model contain in total?

Hint: For `nn.Linear(a,b)`, there are  $a \times b$  weights and  $b$  biases.

☐ 12

☐ 22

☒ 19

☐ 15



### Question 2

0.4 / 0.4 pts

You have the sentence: "I love to [MASK] every morning." and your BERT-based model must predict the missing word. Which statement correctly describes BERT's approach in this scenario?

- ☐ BERT is an autoregressive model that generates tokens one at a time, each conditioned on previous tokens.
- ☒ BERT uses Masked Language Modeling, attending to all surrounding tokens to infer the [MASK].
- ☐ BERT is an autoregressive model that generates tokens one at a time, each conditioned on previous tokens.
- ☐ BERT ignores the context and focuses purely on final-layer hidden states.



### Question 3

0.4 / 0.4 pts

Neural networks are said to be "biologically inspired," yet they differ greatly from real brains. Which statement best explains one key difference?

- ☒ Biological neurons have more complex connectivity and signaling patterns than typical artificial neurons.
- ☐ Biological neurons rely exclusively on matrix multiplication to transmit signals.
- ☐ Artificial neurons use chemical neurotransmitters for activations, unlike real neurons.
- ☐ Artificial networks are only trained with linear activations, whereas biology uses nonlinear signals.



### Question 4

0.4 / 0.4 pts

Some claim that "using only linear activations in multiple layers is not beneficial" because it collapses to a single linear operation. Why is stacking purely linear layers generally ineffective?

- ☒ Multiple linear layers combine into the same transform, providing no additional representation power.
- ☐ Linear layers inherently cause vanishing gradients.
- ☐ Purely linear networks can only handle binary classification.
- ☐ Nonlinear layers slow down training, so fewer are better.



### Question 5

0.4 / 0.4 pts

You have two 300-dimensional word embeddings,  $v_1$  and  $v_2$ . You want to measure how closely they match in meaning. Which measure is most commonly used to assess their similarity?

- ☒ Cosine similarity
- ☐ Euclidean distance

- ☐ Mean squared error
- ☐ Cross-entropy loss



### Question 6

0.4 / 0.4 pts

You discover that your model can be fooled by carefully crafted inputs, causing nonsensical outputs. In LLM terminology, such deceptive or misleading inputs are known as:

Hint: These are maliciously designed prompts or inputs to break or confuse the model

- ☒ Adversarial examples
- ☐ Instruction-tuned prompts
- ☐ Autoregressive prompts
- ☐ Vocabulary collisions



### Question 7

0.4 / 0.4 pts

You are building a neural network to predict housing prices (a continuous value). Which two activation functions are generally suitable at the output layer for a continuous prediction task?

- ☐ Sigmoid
- ☒ Linear
- ☐ Softmax
- ☒ ReLU



### Question 8

0.4 / 0.4 pts

In a typical feed-forward neural network architecture (with an input layer, one or more hidden layers, and an output layer), which of the following is not a usual term for a component in that architecture?

- ☐ Neurons
- ☒ Axons
- ☐ Layers
- ☐ Activation function



### Question 9

0.4 / 0.4 pts

You have a large pre-trained language model and want it to follow specific user instructions (e.g., “turn these notes into a summary”). Which method refers to training a model explicitly with instruction-oriented examples?

Hint: One approach uses tasks posed as explicit instructions (e.g., in a Q&A or summarization format).

- ☐ Increasing the Context Window
- ☐ Fine-Tuning on unstructured text
- ☐ Prompt Engineering
- ☒ Instruction Tuning



IncorrectQuestion 10

0 / 0.4 pts

Assume a feed-forward network with 4 input features, then a single hidden layer of 5 neurons, and an output layer of size 1 for regression. All hidden neurons use ReLU. How many weight parameters (not counting biases) connect the input layer to the hidden layer?

- ☐  $4 \times 5 = 20$
- ☐  $4 + 5 = 9$
- ☒  $4 \times 5 + 5 = 25$
- ☐  $5 \times 1 = 5$

Quiz Score: 3.6 out of 4