

# Particle Filter and Monte Carlo Localization (MCL)

EE468/CE468: Mobile Robotics

---

Dr. Basit Memon

Electrical and Computer Engineering  
Habib University

November 13, 2023



# Table of Contents

- 1 Particle Filter
- 2 How does it work?
- 3 Sources of error in the PF
- 4 References



# Dealing with computational intractability of Bayes

- Markov localization (Bayes Algorithm) is computationally intractable.
  - Closed-form solution seldom exists.

## Parametric Filters (Kalman)

- + Exact filter
- + Computationally efficient
- Exact for linear only
- Inaccurate, when true posterior is multi-modal

## Grid-based

- + Easy - discretize space.
- Curse of dimensionality  
Dynamic discretization
- Wasted computations  
Lookup table
- Wasted memory  
Update select cells

## Monte-Carlo

- Randomly sample space
- + Computationally better.
- + Accommodates arbitrary distribution.
-



# Table of Contents

- 1 Particle Filter
- 2 How does it work?
- 3 Sources of error in the PF
- 4 References



# Key idea of a particle filter:



- Recall **Bayes Filter**

$$\overline{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$$

$$bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$$

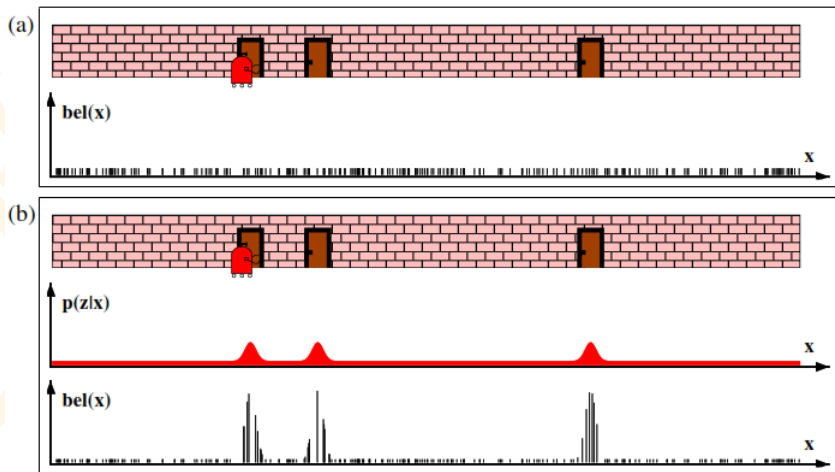
- Represent posterior by set  $\mathcal{M}$  of  $N$  samples (particles).
- Each sample,  $(x^{[i]}, w^{[i]})$ , consists of state hypothesis  $x^{[i]}$  and an importance weight  $w^{[i]}$ .

# Can samples approximate a function?




- The more particles fall in an interval, the higher the probability of that interval.
- $x$  can be repeated in the samples.

# Example of MonteCarlo Localization




# Example of MonteCarlo Localization





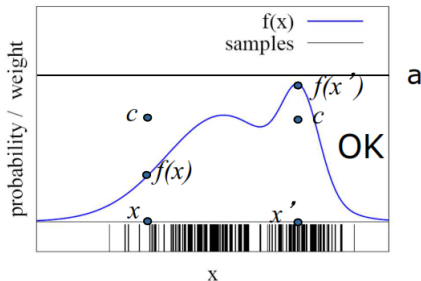
# Example of MonteCarlo Localization





# Table of Contents


- 1 Particle Filter
- 2 How does it work?
- 3 Sources of error in the PF
- 4 References



- Rejection sampling is one method to obtain samples to approximate  $f(x)$ .

- Assume that  $f(x) < a$  for all  $x$ . Then,

- 1 Sample  $x$  from a uniform distribution.
- 2 Sample  $c$  from  $[0, a]$
- 3 If  $f(x) > c$ , then keep the sample, otherwise reject it.



- We currently have

$$bel(x_{t-1}) = \left\{ \begin{array}{l} x_{t-1}^{[1]}, x_{t-1}^{[2]}, \dots, x_{t-1}^{[N]} \\ w_{t-1}^{[1]}, w_{t-1}^{[2]}, \dots, w_{t-1}^{[N]} \end{array} \right\}$$

- Simulate what is going to happen to each particle at next time step given the the input, and select one sample from it as the update for the particle.

$$x_t^{[m]} \sim p(x_t | x_{t-1}^{[m]}, u_t)$$






- How do we sample from

$$bel(x_t) = \eta p(z_t|x_t) \int p(x_t|u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}?$$

- Importance Sampling
  - Sample from an easy proposal distribution
  - Reweigh samples to fix it.



# Importance Sampling: Sample $g$ and reweigh to represent $f$



$$\begin{aligned} E_{f(x)}[X] &= \sum x f(x) \\ &= \sum x f(x) \frac{g(x)}{g(x)} \\ &= \sum x g(x) \frac{f(x)}{g(x)} \\ &= E_{g(x)} \left[ \frac{f(x)}{g(x)} x \right] \end{aligned}$$

- Weigh the samples by  $f(x)/g(x)$ .
- For convergence, wherever  $f(x) > 0$ , the function  $g(x) > 0$ .



- Target distribution:

$$bel(x_t) = \eta p(z_t|x_t) \overline{bel}(x_t)$$

- Proposal distribution:

$$\overline{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$$

- Weighing ratio:

$$\frac{bel(x_t)}{\overline{bel}(x_t)} = \eta p(z_t|x_t)$$

- Weights are proportional to measurement likelihood.



■

$$bel(x_{t-1}) = \left\{ \begin{array}{l} x_{t-1}^{[1]}, x_{t-1}^{[2]}, \dots, x_{t-1}^{[N]} \\ w_{t-1}^{[1]}, w_{t-1}^{[2]}, \dots, w_{t-1}^{[N]} \end{array} \right\}$$

■ for  $i = 1$  to  $N$

$$\text{sample } \bar{x}_t^{[i]} \sim p(x_t | u_t, x_{t-1}^{[i]})$$

$$w_t^{[i]} = \frac{w_t^{[i]}}{\sum w_t^{[i]}}$$

■ for  $i = 1$  to  $N$

$$w_t^{[i]} = p(z_t | \bar{x}_t^{[i]}) w_{t-1}^{[i]}$$

■

$$bel(x_t) = \left\{ \begin{array}{l} \bar{x}_t^{[1]}, \bar{x}_t^{[2]}, \dots, \bar{x}_t^{[N]} \\ w_t^{[1]}, w_t^{[2]}, \dots, w_t^{[N]} \end{array} \right\}$$



## Problem: Unlikely particles consuming compute.

- Waste of computational resources if we propagate particles with too low weights.
- **Solution:** Survival of the fittest! Focus on most likely particles - resample.



# Solution: Resampling

- **Given:** Set  $\mathcal{M}$  of weighted samples
- **Required:**  $N$  samples where the probability of drawing  $x^{[i]}$  is proportional to  $w^{[i]}$ .
- Create a bag of  $x^{[i]}$  from  $\mathcal{M}$ , with instances of  $x^{[i]}$  proportional to  $w^{[i]}$ .
- Draw  $N$  samples from the bag with replacement.
- Weights get reflected in the frequency of particles.



$$bel(x_{t-1}) = \left\{ \begin{array}{c} x_{t-1}^{[1]}, x_{t-1}^{[2]}, \dots, x_{t-1}^{[N]} \\ w_{t-1}^{[1]}, w_{t-1}^{[2]}, \dots, w_{t-1}^{[N]} \end{array} \right\}$$

- for  $i = 1$  to  $N$

$$\text{sample } \bar{x}_t^{[i]} \sim p(x_t | u_t, x_{t-1}^{[i]})$$

$$w_t^{[i]} = \frac{w_t^{[i]}}{\sum w_t^{[i]}}$$


- for  $i = 1$  to  $N$

$$w_t^{[i]} = p(z_t | \bar{x}_t^{[i]}) w_{t-1}^{[i]}$$


- for  $i = 1$  to  $N$ , sample  $x_t^{[i]} \sim w_t^{[i]}$

$$bel(x_t) = \left\{ \begin{array}{c} x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[N]} \\ 1, 1, \dots, 1 \end{array} \right\}$$

# Resampling



- Roulette wheel
- Binary search,  $n \log n$



- Stochastic universal sampling
- Systematic resampling
- Linear time complexity
- Easy to implement, low variance

Slide Credit: Probabilistic Robotics



# Roulette Sampling or Fitness Proportionate Selection

- $\eta = w^1 + w^2 + \dots + w^N.$

- Create cumulative weights array  $c$ , such that

$$c_i = (w^1 + w^2 + \dots w^i)/\eta.$$

- Draw  $N$  random numbers  $u_1, \dots, u_N \sim U[0, 1]$ .
- For each  $j = 1$  to  $N$ , find the first normalized sum entry  $i$  such that  $u_j < c_i$ .
- Add sample  $x^{i-1}$  to new particle set.

# Resampling Algorithm

1. Algorithm **systematic\_resampling**( $S, n$ ):
2.  $S' = \emptyset, c_1 = w^1$
3. **For**  $i = 2 \dots n$  *Generate cdf*
4.  $c_i = c_{i-1} + w^i$
5.  $u_1 \sim U[0, n^{-1}]$ ,  $i = 1$  *Initialize threshold*
6. **For**  $j = 1 \dots n$  *Draw samples ...*
7. **While** ( $u_j > c_i$ ) *Skip until next threshold reached*
8.  $i = i + 1$
9.  $S' = S' \cup \{x^i, n^{-1}\}$  *Insert*
10.  $u_{j+1} = u_j + n^{-1}$  *Increment threshold*
11. **Return**  $S'$

Also called **stochastic universal sampling**



# Table of Contents

- 1 Particle Filter
- 2 How does it work?
- 3 Sources of error in the PF**
- 4 References







# Problem: Sampling Variance




- Whenever a finite number of samples are drawn from any PDF, the statistics of the samples will differ from the underlying PDF.
- **Sampling variance is amplified through repetitive resampling.**
- Suppose a robot does not move, but we don't know its location. Prior is uniform and PF starts with particles spread out. During resampling, PF will occasionally fail to reproduce a sample,  $x^{[m]}$ . Over time, with probability 1 we'll be left with identical copies of a single particle.



(a) 25 samples



(b) 250 samples






# Sampling Variance can be reduced by low-variance sampling.

- Stochastic universal sampling is a low-variance sampling method.

# Problem: Particle Deprivation



- Particles can get stuck in regions of low probability. No particles in the vicinity of correct state.
- Occurs mostly when not enough particles.
- Could happen because of sampling variance or an unlucky series of random numbers wipe out particles near true state.



## Solution: Injection of random particles

- Robot might get kidnapped with a small probability at any time step.
- How many particles? From which distribution?

```

1:  Algorithm Augmented_MCL( $\mathcal{X}_{t-1}, u_t, z_t, m$ ):
2:      static  $w_{\text{slow}}, w_{\text{fast}}$ 
3:       $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
4:      for  $m = 1$  to  $M$  do
5:           $x_t^{[m]} = \text{sample\_motion\_model}(u_t, x_{t-1}^{[m]})$ 
6:           $w_t^{[m]} = \text{measurement\_model}(z_t, x_t^{[m]}, m)$ 
7:           $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
8:           $w_{\text{avg}} = w_{\text{avg}} + \frac{1}{M} w_t^{[m]}$ 
9:      endfor
10:      $w_{\text{slow}} = w_{\text{slow}} + \alpha_{\text{slow}}(w_{\text{avg}} - w_{\text{slow}})$ 
11:      $w_{\text{fast}} = w_{\text{fast}} + \alpha_{\text{fast}}(w_{\text{avg}} - w_{\text{fast}})$ 
12:     for  $m = 1$  to  $M$  do
13:         with probability  $\max\{0.0, 1.0 - w_{\text{fast}}/w_{\text{slow}}\}$  do
14:             add random pose to  $\mathcal{X}_t$ 
15:         else
16:             draw  $i \in \{1, \dots, N\}$  with probability  $\propto w_t^{[i]}$ 
17:             add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
18:         endwith
19:     endfor
20:     return  $\mathcal{X}_t$ 

```

**Table 8.3** An adaptive variant of MCL that adds random samples. The number of random samples is determined by comparing the short-term with the long-term likelihood of sensor measurements.



# Table of Contents

- 1 Particle Filter
- 2 How does it work?
- 3 Sources of error in the PF
- 4 References

